



KE UNIT 3 DATA WAREHOUSING FOR BUSINESS ANALYTICS

DAY 1

Dr TIAN Jing

tianjing@nus.edu.sg



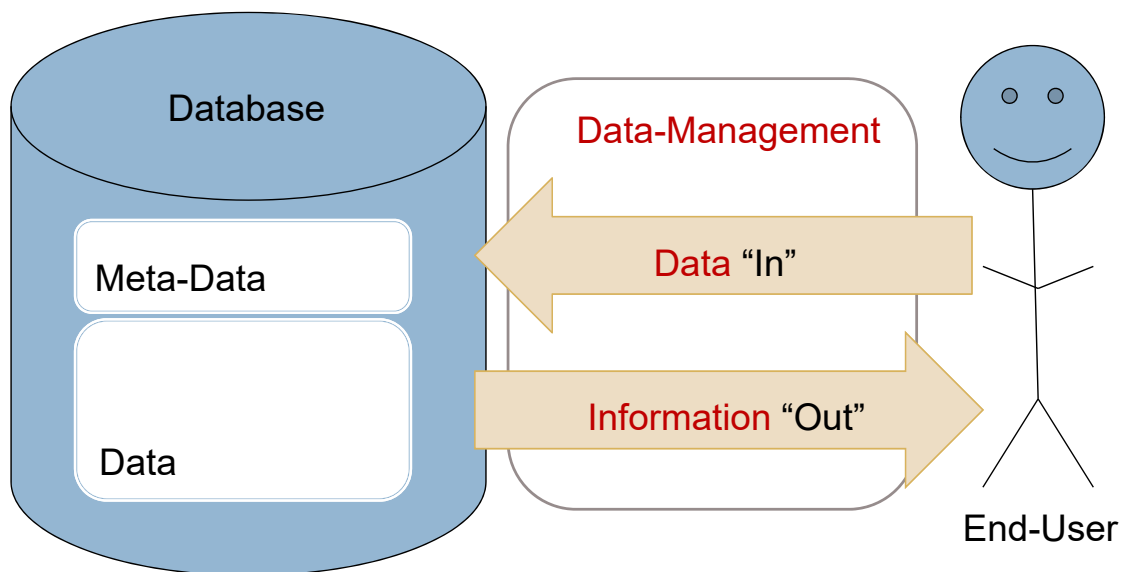
Lesson plan (part A)

| Monday 16 July | Tuesday 17 July | Wednesday 18 July | Thursday 19 July | Friday 20 July |
|---|--|---|---|---|
| 9.00 - 12.00 | 9.00 - 12.00 | 9.00 - 12.00 | 9.00 - 12.00 | 9.00 - 12.00 |
| 1b. Introduction to Data Modelling (I) | 2a. Relational Database and SQL (I) | 3a. Introduction to Data Warehousing | 4a. Data Visualisation and Storytelling (II) | 5a. Project Consultation and Presentation |
| Tian Jing | Tian Jing | Brandon | Brandon | Brandon / Tian Jing |
| 13.30 - 17.00 | 13.30 - 17.00 | 13.30 - 17.00 | 13.30 - 17.00 | 13.30 - 17.00 |
| 1b. Introduction to Data Modelling (II) | 2b. Relational Database and SQL (II) | 3b. Data Visualisation and Storytelling (I) | 4b. Data Visualisation and Storytelling (III) | 5b. Project Consultation and Presentation |
| Tian Jing | Tian Jing | Brandon | Brandon | Brandon / Tian Jing |

- Team: 4-5 students
- CA1 project presentation (5%)
 - 10 minutes per team
 - 5th day 20 July (Friday)
 - Concept and architecture design ONLY
- CA1 submission (15%)
 - Due date: 12 August (Sunday)
 - Refer to CA1 briefing document in IVLE
- Exam (part A): 30%

- Introduction to data modelling
- Relational data modelling
 - Entity relationship diagram (ERD).
 - Normalization, a technique that helps analysts validate the data models
 - Logical data modelling
 - Data management and query (Day 2)
- Dimensional data modelling (Day 3)
- Non-relational data modelling (part B)

- Introduction to data modelling
- Data analysis
 - Entity relationship diagram (ERD)
 - Attribute analysis
- Data design
 - Normalization
 - Logical data model



| | |
|------------------------|---|
| Data | Data are raw unprocessed facts. By itself data has no meaning and no structure. |
| Information | Information is interpreted or processed data. |
| Data Management | <ul style="list-style-type: none"> •Create - adding new data •Read - retrieving information •Update - modifying existing data •Delete - removing data |
| Metadata | Data structure and category |
| Query | Asking questions of data in search of a specific answer. |

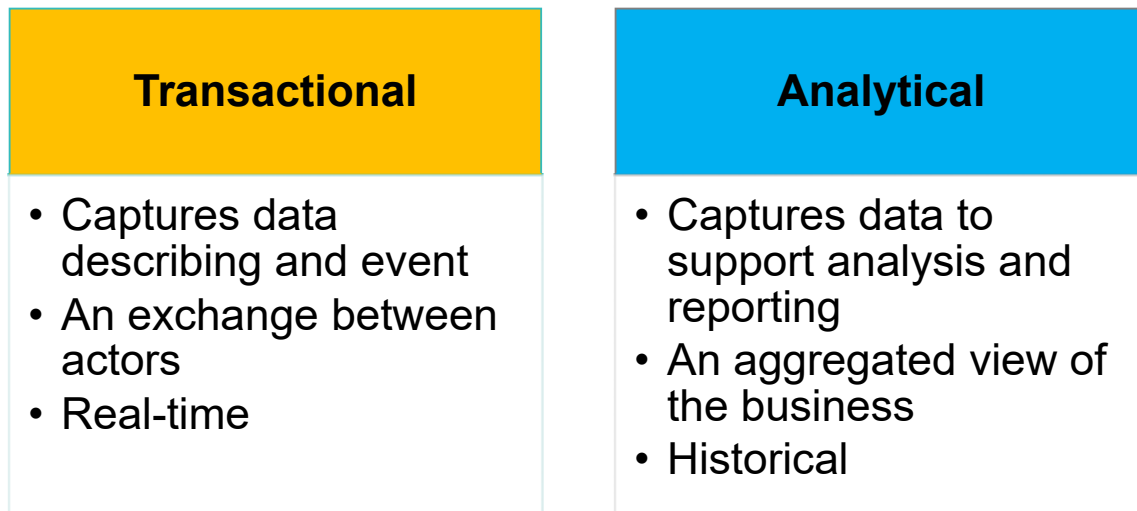


Activity

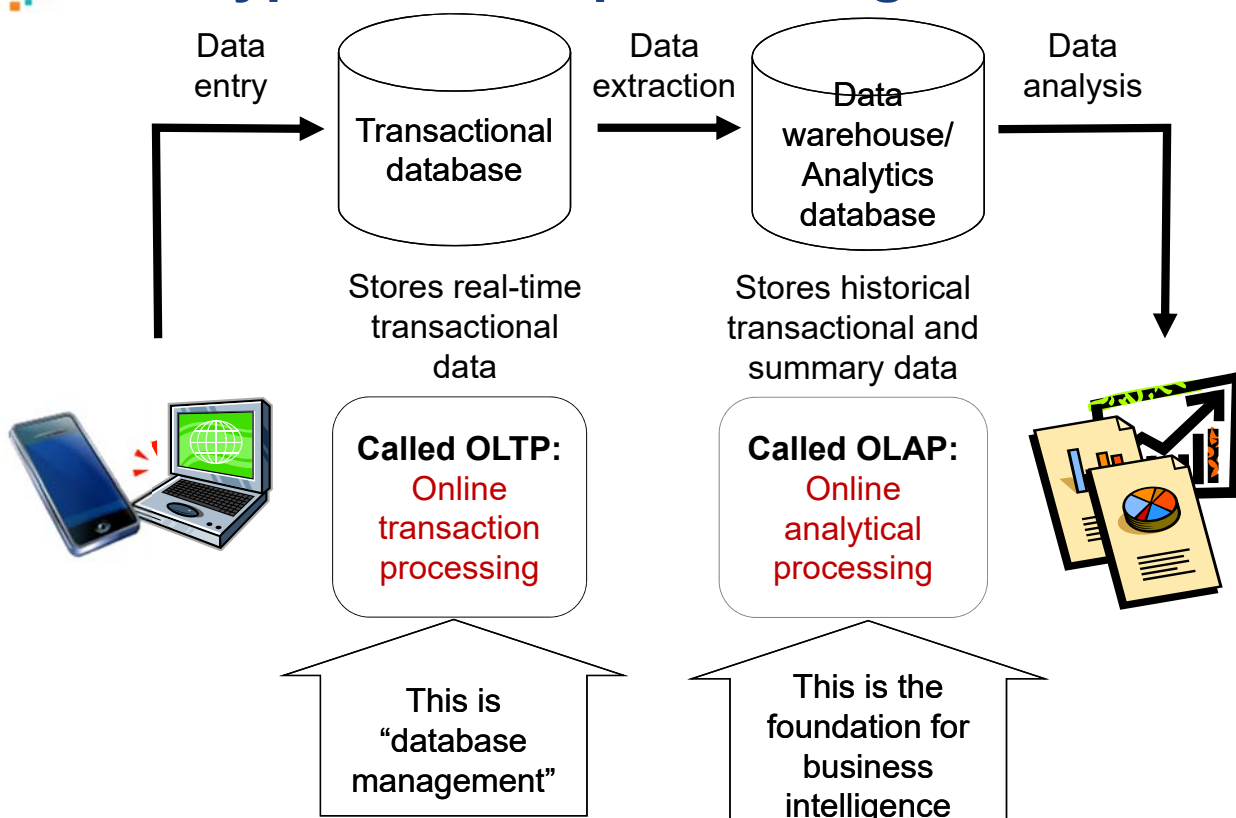
1. A Telephone book
2. Organizing the Phone Book in Alphabetical Order
3. Looking up 'Michael Fudge' yields the phone number 555-1234
4. How many 'Fudges' are there in the phone book?
5. Employee records (in a file cabinet)
6. Filing a new employee under "W" because their last name is "Williams"
7. The average employee salary is \$40,000

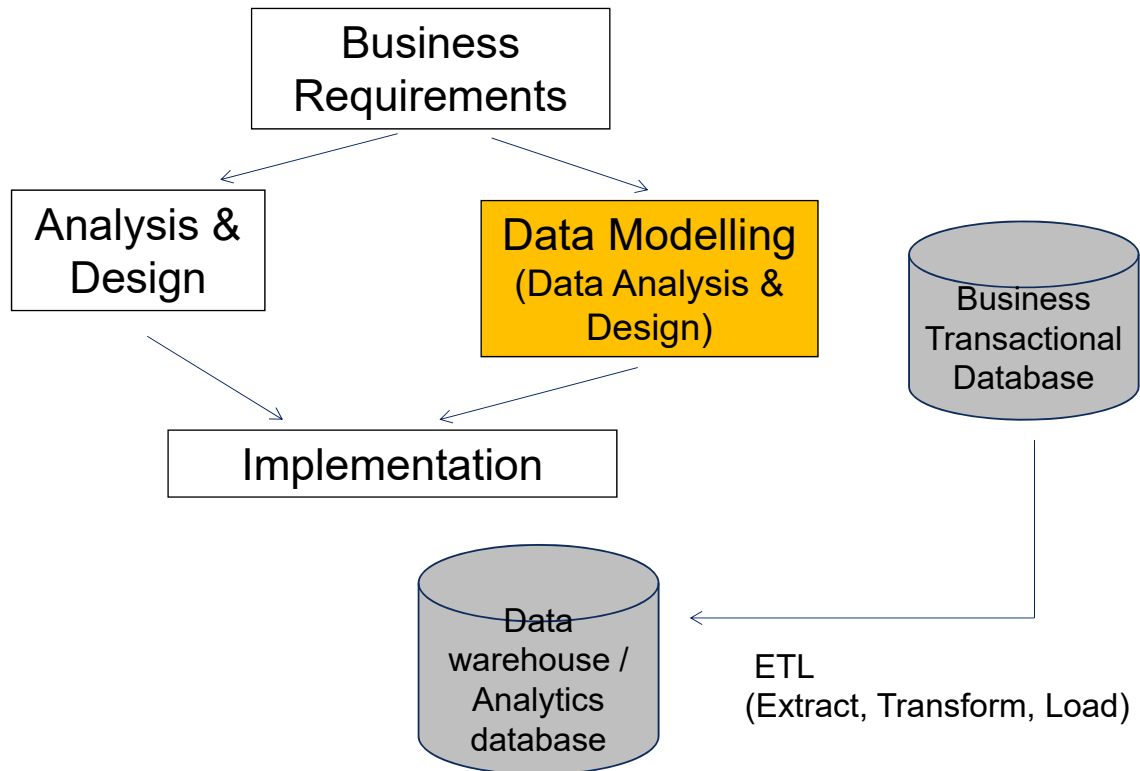
| | |
|------------------|--|
| Data? | |
| Information? | |
| Data Management? | |
| Metadata? | |
| Query? | |

Two types of data



Two types of data processing



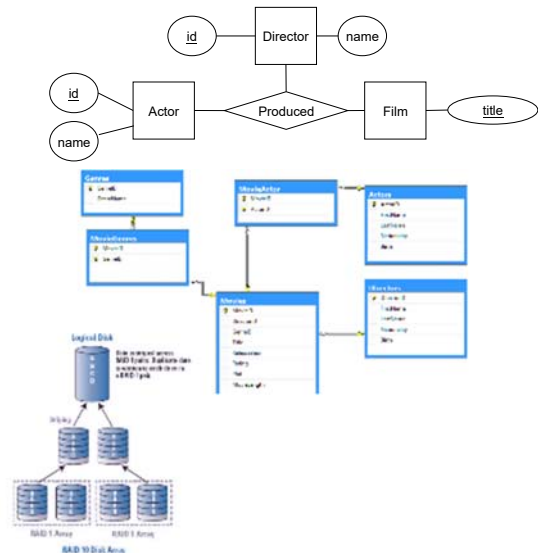


Data modelling

- What is a **database**?
 - A collection of files storing related data
- What is a **database management system** (DBMS)?
 - An application program that allows us to manage efficiently the collection of data files.
- What is **data model**?
 - Mathematical formalism (or conceptual way) for describing the data


Data modelling

- **Requirement analysis**
 - What information needs to be stored? How will it be used? What integrity constraints should be imposed?
- **Conceptual data modelling**
 - Define/describe/discuss the semantic modeling of data in the application (Entities Relations Diagrams)
- **Logical data modelling**
 - Enhance the Entities Relations Diagrams by optimizations and relationships
- **Physical data modelling**
 - Translate the database schema into a physical storage plan on available hardware (DBMS, SQL, day 2)



Data analysis and data design

- **Data Analysis**
 - Develop a data model for data required by the business requirements
 - A data model consists of
 - Entities Relations Diagram
 - Data Dictionary
- **Data Design**
 - Restructure the data model so that it is optimised or suitable for data accesses
 - Required before designing and implementing physical database
 - Logical data model (through normalization)

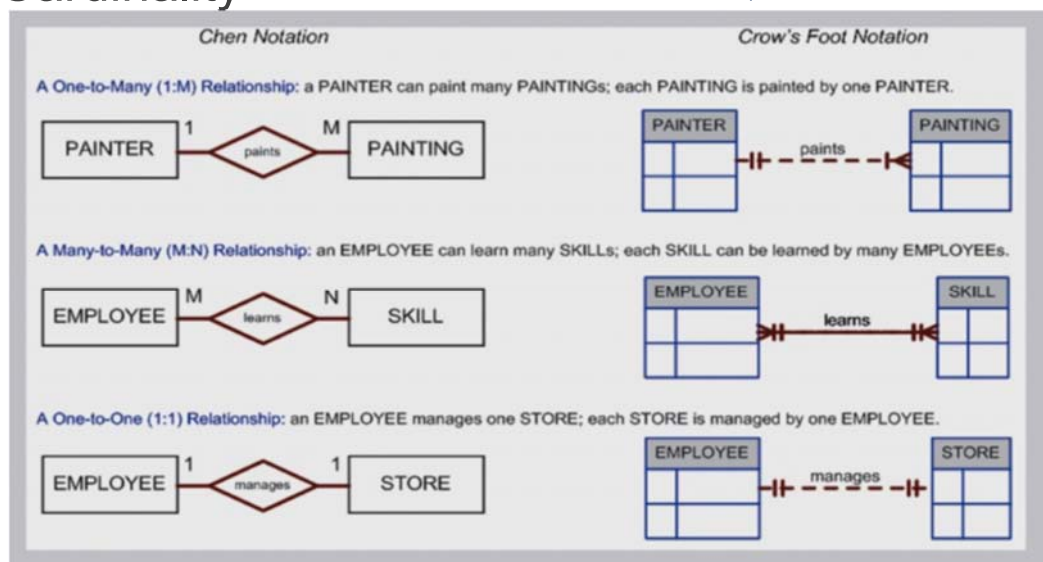


- Introduction to data modelling
- Data analysis
 - Entity relationship diagram (ERD)
 - Attribute analysis
- Data design
 - Normalization
 - Logical data model

Entity relationship diagram (ERD)

- Entities
- Relationships
- Cardinality


We use **Crow's Foot Notation** in our class



- A distinguishable objects in the problem domain that we want to model.
- You need to distinguish:
 - Entity Type (or Entity)
 - Entity Occurrence

Example: In a ISS Course Registration System


- Teacher is an Entity Type
- Brandon and I are the Entity Occurrences



- Depends on the business rules
- Every Relationship is bi-directional
 - a teacher teaches one or more students
 - a student is taught by one or more teacher



- There may be more than one important relationships
 - a teacher counsels zero, one or more students
 - a teacher teaches one or more students
- Every Relationship is described in terms of a verb



- Cardinality of relationships: How many
 - The number of occurrences of one entity type that relate to the occurrences of another entity type


| | | | |
|---------------|---|---|---|
| one- to- one | + | — | + |
| one- to- many | + | — | < |
| many-to-many | > | — | < |

- Optional Relationships

| | | | |
|------------------------|---|---|----|
| • one to zero or one | + | — | ○+ |
| • one to zero or many | + | — | ○< |
| • many to zero or many | > | — | ○< |

Example (1)


- An education institute management system



Example (2)

- Data modelling is similar to diagramming a sentence
 - Place boxes around the 'nouns', or entities.
 - Underline the 'verbs'.
 - Circle the 'how many' qualifier.
 - Look for optionality words such as 'may/must'.

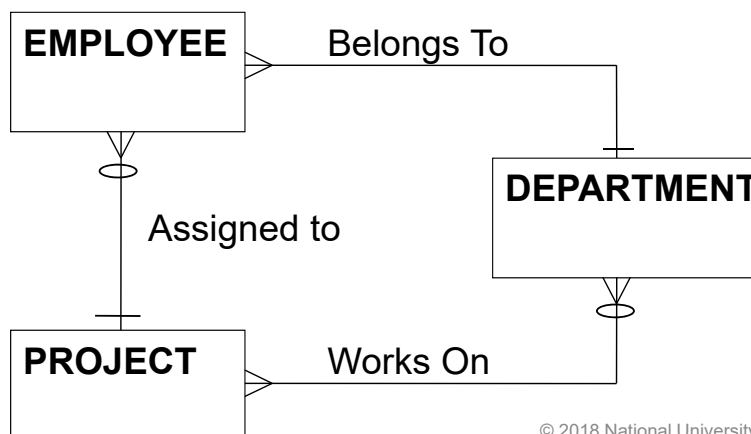
A **department** may contain more than one **employee**.



Example (3)

A project management system with objective to keep track of

- all the projects undertaken by the company
- assignments of projects to departments
 - a project can be assigned to one or more departments
 - each department can take on one or more projects
 - one or more employee of a department may work on a project
 - a employee can only take on one project at any one time



- Introduction to data modelling
- Data analysis
 - Entity relationship diagram (ERD)
 - **Attribute analysis**
- Data design
 - Normalization
 - Logical data model

- Meta Data
 - Data about data
- Data Dictionary
 - a repository to store all information, description about the data
 - Contains metadata

| ATTRIBUTE NAME | TYPE | LENGTH | DEFINITIONS AND BUSINESS RULES |
|------------------|------|--------|---|
| EMP-M | A | 25 | Employee Name (full name, start with the surname) |
| Emp-JOB-T | A | 25 | Employee Job title (Programmer, Analyst, Project Manager, Department Manager) |
| EMP-JOB-DESC | A | 60 | Simple short description |
| PROJ-M | A | 10 | A unique short name given to the project |
| MTH-SAL-A | N | 6.2 | Monthly salary (999999.99) |
| EMP-PROJ-START-D | D | 8 | Start date of the project (DDMMYYYY) |
| EMP-PROJ-END-D | D | 8 | End date of the project (DDMMYYYY) |

Note: A – Alphanumeric, N – Numeric, D - date



Example


- Recall for the project management system, possible attributes of each entity are
 - Employee
 - employee name, job title, job description, project name, employee skill type, employee skill type description, employee monthly salary etc
 - Department
 - department name, department manager number, department manager name, department employee size, project name, project-department budget allocated, department employee number, department employee name
 - Project
 - project name, project description, project budget allocated, project start date, project end date

- Attributes of the Project Management system

| ENTITY | ATTRIBUTE NAME | TYPE | LEN | DEFINITION AND BUSINESS RULES |
|----------|------------------|------|-----|--|
| Employee | EMP-N | N | 6 | Employee number. Unique for each employee. |
| | EMP-M | A | 25 | Employee name in the form of the last name and two initials. |
| | EMP-JOB-T | A | 25 | Employee job Title (analyst, programmer, project manager, department manager) |
| | EMP-JOB-DESC | A | 60 | Short description of job title |
| | SAL-CHNG-D | N | 8 | The effective date of the employee's salary.. Employee salary history is kept for 3 years. |
| | MTH-SAL-A | N | 6.2 | Monthly salary (in the form of 999999.99). |
| | PROJ-M | A | 10 | Name of project currently assigned to employee Each project has a unique name |
| | EMP-PROJ-START-D | D | 8 | Start date of the employee on the project |
| | EMP-PROJ-END-D | D | 8 | End date of the employee on the project |
| | SKILL-TYPE-C | A | 6 | Skill type code |
| | SKILL-TYPE-DESC | A | 20 | Description of the skill (usually an employee has more than one skill) |

| ENTITY | ATTRIBUTE NAME | TYPE | LEN | DEFINITION AND BUSINESS RULES |
|------------|-------------------|------|-----|--|
| Department | DEPT-M | A | 10 | A unique name for each department |
| | DEPT-MGR-N | N | 6 | Employee number of the manager in charge of department |
| | DEPT-MGR-M | A | 25 | Name of the manager in charge of department |
| | DEPT-EMP-SIZE-Q | N | 3 | Number of employee in the department |
| | DEPT-PROJ-M | A | 10 | Project (name) currently assigned to department |
| | DEPT-PROJ-BUDGT-A | N | 6.2 | Allocated budget(money) for the project for that department |
| | DEPT-EMP-N | N | 6 | Employee Numbers of all employees working in the department |
| | DEPT-EMP-M | A | 25 | Employee names of all employees working in the department |
| Project | PROJ-M | A | 10 | Unique name assigned to project. Project currently work on by company. |
| | PROJ-DESC | A | 100 | Short description of the project |
| | PROJ-BUDGT-A | N | 6.2 | Total budget (money) allocated to the entire project |
| | PROJ-START-D | D | 8 | Project start date (DDMMYYYY) |
| | PROJ-END-D | D | 8 | Project end date (DDMMYYY). The project must end by this date. |

- **Primary key (PK):** unique identifier of a record
 - A simple key is a key consisting of a single attribute
 - A composite key is a key consisting of more than one attribute
 - No two rows must contain the same value for their primary keys
 - None of the component attributes of the identifier may have null values.



Example: Identify key attributes

| ENTITY | ATTRIBUTE NAME | TYPE | LEN | DEFINITION AND BUSINESS RULES |
|----------|------------------|------|-----|--|
| Employee | <u>EMP-N</u> | N | 6 | Employee number. Unique for each employee. |
| | EMP-M | A | 25 | Employee name in the form of the last name and two initials. |
| | EMP-JOB-T | A | 25 | Employee job Title (analyst, programmer, project manager, department manager) |
| | EMP-JOB-DESC | A | 60 | Short description of job title |
| | SAL-CHNG-D | D | 8 | Date an employee's salary was changed, in the form of 'DDMMYYYY'. Employee salary history is kept for 3 years. |
| | MTH-SAL-A | N | 6.2 | Monthly salary after change on a given date (in the form of 999999.99). |
| | PROJ-M | A | 10 | Name of project assigned to employee Each project has a unique name |
| | EMP-PROJ-START-D | D | 8 | Start date of employee assignment to the project, DDMMYYYY (assignment period of employee to the project varies between employees) |
| | EMP-PROJ-END-D | D | 8 | End date of employee assignment to the project, DDMMYYYY (assignment period of employee to the project varies between employees) |
| | SKILL-TYPE-C | A | 6 | Skill type code |
| | SKILL-TYPE-DESC | A | 20 | Description of the skill |



Example: Identify key attributes

| ENTITY | ATTRIBUTE NAME | TYPE | LEN | DEFINITION AND BUSINESS RULES |
|------------|-------------------|------|-----|---|
| Department | <u>DEPT-M</u> | A | 10 | A unique name for each department |
| | DEPT-MGR-N | N | 6 | Employee number of the manager in charge of department |
| | DEPT-MGR-M | A | 25 | Name of the manager in charge of department |
| | DEPT-EMP-SIZE-Q | N | 3 | Number of employee in the department |
| | DEPT-PROJ-M | A | 10 | Project (name) assigned to department |
| | DEPT-PROJ-BUDGT-A | N | 6.2 | Allocated budget(money) for the project for that department |
| | DEPT-EMP-N | N | 6 | Employee Numbers of all employees working in the department |
| | DEPT-EMP-M | A | 25 | Employee names of all employees working in the department |
| Project | <u>PROJ-M</u> | A | 10 | Unique name assigned to project |
| | PROJ-DESC | A | 100 | Short description of the project |
| | PROJ-BUDGT-A | N | 6.2 | Total budget (money) allocated to the entire project |
| | PROJ-START-D | D | 8 | Project start date (DDMMYYYY) |
| | PROJ-END-D | D | 8 | Project end date (DDMMYYYY). The project must end by this date. |



Summary: Design ERD

- **Identify the entities**
 - If you begin the data model using a use case, look at the major inputs and outputs of the use case.
 - If the process models are available, look at the data stores, external entities, and data flows.
- **Add the appropriate attributes to each entity**
 - One or more of the attributes will become the entity's identifier.
- **Draw relationships among entities**
 - Each relationship is labeled, and cardinality and modality are assigned.

- Introduction to data modelling
- Data analysis
 - Entity relationship diagram (ERD)
 - Attribute analysis
- Data design
 - **Normalization**
 - Logical data model



What is normalization?

- A technique to organize **“efficiently”** organize data in a database
- **“Efficiently”**:
 - Eliminating redundant data
 - Not storing the same data in more than one table
 - Ensuring that functional dependencies make sense

Without normalization

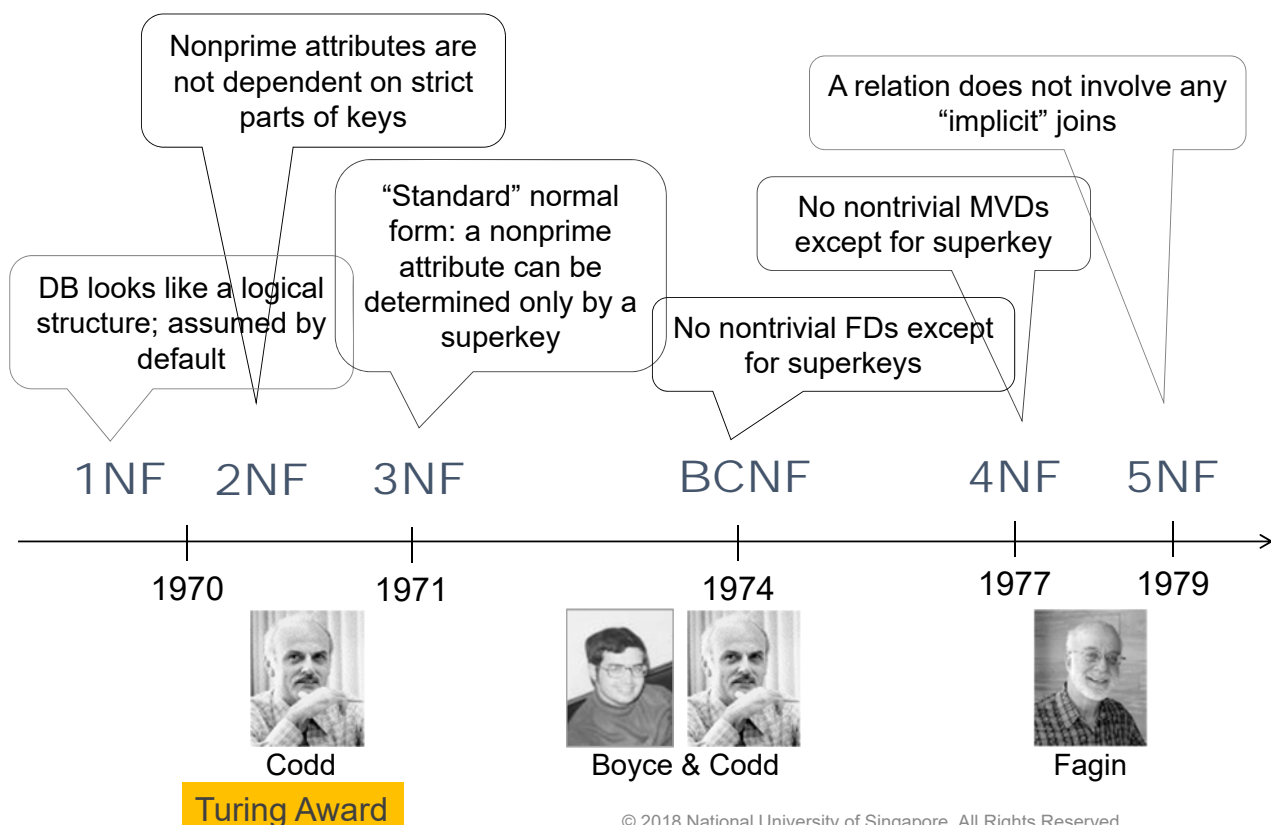
| student_id | name | address | subject |
|------------|--------|----------------|---------|
| 401 | Adam | 133 Our Lane | Biology |
| 402 | Alex | 123 Here Lane | Math |
| 403 | Stuart | 123 My Lane | Math |
| 404 | Adam | 123 Their Lane | Physics |

•**Update Anomaly** : To update address of a student who occurs twice or more than twice in a table, we will have to update **address** column in all the rows, else data will become inconsistent.

•**Insertion Anomaly** : Suppose for a new admission, we have a Student id(S_id), name and address of a student but if student has not opted for any subjects yet then we have to insert **NULL** there, leading to Insertion Anomaly.

•**Deletion Anomaly** : If (student_id) 401 has only one subject and temporarily he drops it, when we delete that row, entire student record will be deleted along with it.

History of normal forms





Business case: Order processing

Id: A1091

Purchase Order

Customer ID: *S009*
Customer Name: *Lynn Wang*
Date: *21/7/2011*

| ItemCode | Description | Qty |
|----------|-------------|-----|
| S1001 | Pencil | 100 |
| S1003 | Eraser | 200 |
| S1005 | Ruler | 250 |


Total Number of Items: 3



Product Code : S1001
Product Name: Pencil
Unit Price: \$0.20

Business Rules:


- An order contain 1 to many products
- A product may be appear in multiple order
- A product may not be ordered if it is not popular!



Initial design of ERD

| Orders |
|--------------------|
| <u>OrderID</u> |
| CustomerID |
| CustomerName |
| OrderDate |
| ProductID |
| ProductName |
| Qty |
| ProductsTotal |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |



| | |
|--|--|
| ONF → 1NF | <ul style="list-style-type: none"> Multivalue Attributes/Repeating Groups? <ul style="list-style-type: none"> Move multivalue/repeating group to a new entity -determine the key of the new relation |
| 1NF → 2NF | <ul style="list-style-type: none"> Are there attributes dependent on a partial key (of composite key)? <ul style="list-style-type: none"> Move attribute(s) to a new entity -determine the key of the new entity |
| 2NF → 3NF | <ul style="list-style-type: none"> Any non-key attribute dependent on any other non-key attribute? <ul style="list-style-type: none"> Move attribute(s) to a new entity -determine the key of the new entity |
| Optimization - combining entities with same primary key / remove derivable attribute | |

Question

Question: Check all tables with multivalue repeating attributes?

| OrderID [PK] | CustomerID | Customer Name | Order Date | ProductID | Product Name | Qty | Products Total |
|--------------|------------|---------------|------------|-----------|--------------|-----|----------------|
| A1091 | S009 | Lynn Wang | 21/7/2011 | S1001 | Pencil | 100 | 3 |
| | | | | S1003 | Eraser | 200 | |
| | | | | S1005 | Ruler | 250 | |
| A1092 | S010 | Suzan Tan | 12/1/2011 | S1001 | Pencil | 10 | 2 |
| | | | | S1004 | Pen | 50 | |
| A1093 | S010 | Suzan Tan | 21/8/2011 | S1001 | Pencil | 80 | 2 |
| | | | | S1005 | Ruler | 90 | |

Multivalue attribute

Multivalue attribute

Multivalue attribute

Repeating group
(group of related multivalue attributes)

ONF → 1NF: Action required

- If there are multivalued attributes and/or repeating groups
 - **Place** each attribute/group into a separate new table
 - **Copy** the primary key from the original table to the new tables
- Examine the new table and **determine** which additional attribute(s) are needed to uniquely identify a single row of the new table. The primary key from the original table usually is insufficient to be the primary key in the new table
- **Give** names to the new table

Example

Orders

| OrderID [PK] | CustomerID | CustomerName | OrderDate | ProductsTotal |
|--------------|------------|--------------|-----------|---------------|
| A1091 | S009 | Lynn Wang | 21/7/2011 | 3 |
| A1092 | S010 | Suzan Tan | 12/1/2011 | 2 |
| A1093 | S010 | Suzan Tan | 21/8/2011 | 2 |

OrderDetails

| OrderID [PK] | ProductID [PK] | Product Name | Qty |
|--------------|----------------|--------------|-----|
| A1091 | S1001 | Pencil | 100 |
| A1091 | S1003 | Eraser | 200 |
| A1091 | S1005 | Ruler | 250 |
| A1092 | S1001 | Pencil | 10 |
| A1092 | S1004 | Pen | 50 |
| A1093 | S1001 | Pencil | 80 |
| A1093 | S1005 | Ruler | 90 |

0NF → 1NF

| Orders |
|--------------------|
| <u>OrderID</u> |
| CustomerID |
| CustomerName |
| OrderDate |
| ProductID |
| ProductName |
| Qty |
| ProductsTotal |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |

0NF (top) → 1NF (bottom)
We have split table(s) with multi-value attributes or repeating group

| Orders |
|----------------|
| <u>OrderID</u> |
| CustomerID |
| CustomerName |
| OrderDate |
| ProductsTotal |

| OrderDetails |
|------------------|
| <u>OrderID</u> |
| <u>ProductID</u> |
| ProductName |
| Qty |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |

Next step (1NF → 2NF)

0NF → 1NF

• Multivalue Attributes/Repeating Groups?

- Move multivalue/repeating group to a new entity
- determine the key of the new relation

1NF → 2NF

• Are there attributes dependent on a partial key (of **composite key**)?

- Move attribute(s) to a new entity
- determine the key of the new entity

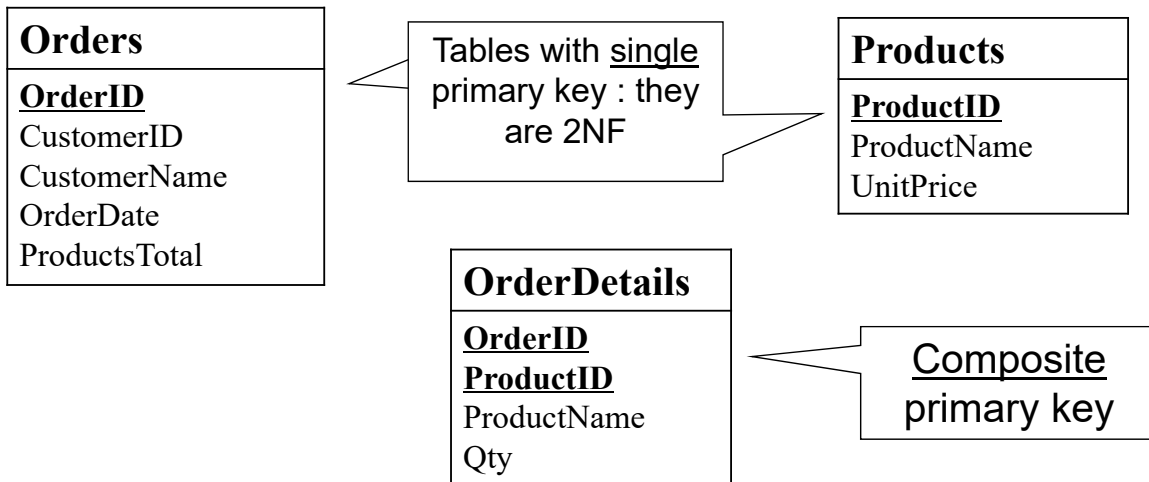
2NF → 3NF

• Any non-key attribute dependent on any other non-key attribute?

- Move attribute(s) to a new entity
- determine the key of the new entity

Optimization - combining entities with same primary key / remove derivable attribute

Question: Check all tables with composite primary key: any attribute depends on part of whole-key?



OrderDetails

Table with composite primary key

| OrderID [PK] | ProductID [PK] | ProductName | Qty |
|--------------|----------------|-------------|-----|
| A1091 | S1001 | Pencil | 100 |
| A1091 | S1003 | Eraser | 200 |
| A1091 | S1005 | Ruler | 250 |
| A1092 | S1001 | Pencil | 10 |
| A1092 | S1004 | Pen | 50 |
| A1093 | S1001 | Pencil | 80 |
| A1093 | S1005 | Ruler | 90 |

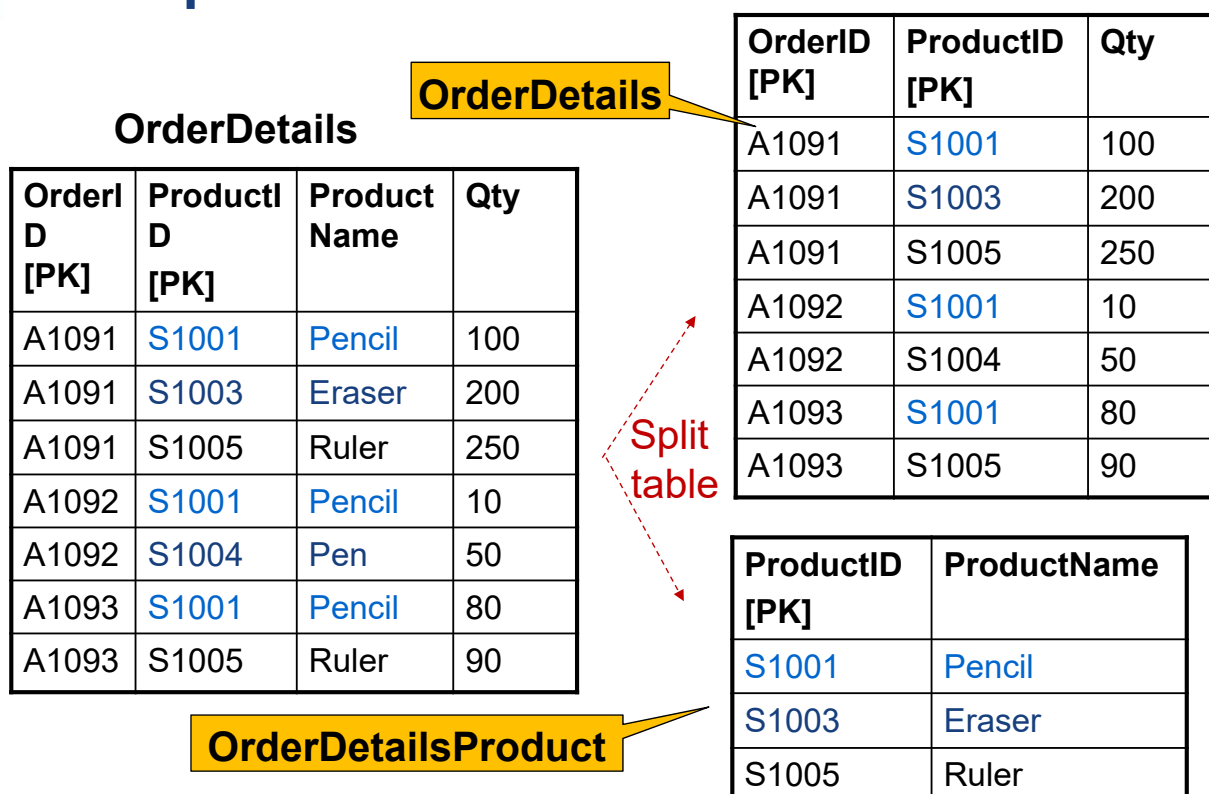
↑

ProductName depends on ProductID (part of composite key) and not OrderID

1NF → 2NF: Action required

- Table with a single primary key is already 2NF
- Table with a compose primary key
 - **Check** each attribute against the whole key, **move** attribute(s) and the part of the key on which it depends to form a new table
 - **Name** the new tables(s)
 - **Decide** on the primary key of the new table

Example



1NF → 2NF

| Orders |
|----------------|
| <u>OrderID</u> |
| CustomerID |
| CustomerName |
| OrderDate |
| ProductsTotal |

| OrderDetails |
|------------------|
| <u>OrderID</u> |
| <u>ProductID</u> |
| ProductName |
| Qty |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |

1NF (top) → 2NF (bottom)

We have split table(s) with attributes dependent on a partial key (of composite key)

| Orders |
|----------------|
| <u>OrderID</u> |
| CustomerID |
| CustomerName |
| OrderDate |
| ProductsTotal |

| OrderDetails |
|------------------|
| <u>OrderID</u> |
| <u>ProductID</u> |
| Qty |

| OrderDetailsProduct |
|---------------------|
| <u>ProductID</u> |
| ProductName |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |

Next step (2NF → 3NF)

ONF → 1NF

• **Multivalue Attributes/Repeating Groups?**

- Move multivalue/repeating group to a new entity
- determine the key of the new relation

1NF → 2NF

• **Are there attributes dependent on a partial key (of composite key)?**

- Move attribute(s) to a new entity
- determine the key of the new entity

2NF → 3NF

• **Any non-key attribute dependent on any other non-key attribute?**

- Move attribute(s) to a new entity
- determine the key of the new entity

Optimization - combining entities with same primary key / remove derivable attribute

- Check all tables: any attribute depends on non-key attribute

Orders

| OrderID [PK] | CustomerID | CustomerName | OrderDate | ProductsTotal |
|--------------|------------|--------------|-----------|---------------|
| A1091 | S009 | Lynn Wang | 21/7/2011 | 3 |
| A1092 | S010 | Suzan Tan | 12/1/2011 | 2 |
| A1093 | S010 | Suzan Tan | 21/8/2011 | 2 |



CustomerName depends on CustomerID (non-key attribute)



2NF → 3NF: Action required

- **Examine** each attribute
- If an attribute(s) does not depend on the whole key, or it depends on another non-key attribute, remove the attribute(s) and use attribute on which it depends on to form a new relation. i.e. **create** new table comprising the attribute(s) and the non-key attribute upon which it depends
- **Determine** the key(s) for the new table(s)
- **Name** the new table(s)

Orders

| OrderID[PK] | CustomerID | CustomerName | OrderDate | ProductsTotal |
|-------------|------------|--------------|-----------|---------------|
| A1091 | S009 | Lynn Wang | 21/7/2011 | 3 |
| A1092 | S010 | Suzan Tan | 12/1/2011 | 2 |
| A1093 | S010 | Suzan Tan | 21/8/2011 | 2 |

Orders

split table

| Order ID [PK] | Customer ID | Order Date | Products Total |
|---------------|-------------|------------|----------------|
| A1091 | S009 | 21/7/2011 | 3 |
| A1092 | S010 | 12/1/2011 | 2 |
| A1093 | S010 | 21/8/2011 | 2 |

Customers

| CustomerID | Customer Name |
|------------|---------------|
| S009 | Lynn Wang |
| S010 | Suzan Tan |



2NF → 3NF

| Orders |
|----------------|
| <u>OrderID</u> |
| CustomerID |
| CustomerName |
| OrderDate |
| ProductsTotal |

| OrderDetails |
|------------------|
| <u>OrderID</u> |
| <u>ProductID</u> |
| Qty |

| OrderDetailsProduct |
|---------------------|
| <u>ProductID</u> |
| ProductName |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |

2NF (top) → 3NF (bottom)

We have split table(s) with any attribute depends on non-key attribute

| Orders |
|----------------|
| <u>OrderID</u> |
| CustomerID |
| OrderDate |
| ProductsTotal |

| Customers |
|-------------------|
| <u>CustomerID</u> |
| CustomerName |

| OrderDetails |
|------------------|
| <u>OrderID</u> |
| <u>ProductID</u> |
| Qty |

| OrderDetailsProduct |
|---------------------|
| <u>ProductID</u> |
| ProductName |

| Products |
|------------------|
| <u>ProductID</u> |
| ProductName |
| UnitPrice |

| Orders | Customers | OrderDetails | OrderDetailsP roduct | Products |
|--|-----------------------------------|---|---------------------------------|--|
| <u>OrderID</u> CustomerID OrderDate ProductsTotal | <u>CustomerID</u> CustomerName | <u>OrderID</u> <u>ProductID</u> Qty | <u>ProductID</u> ProductName | <u>ProductID</u> ProductName UnitPrice |

Optimization: Combine tables with same primary key

| Orders | Customers | OrderDetails | OrderDetailsP roduct | Products |
|--|-----------------------------------|---|---------------------------------|--|
| <u>OrderID</u> CustomerID OrderDate ProductsTotal | <u>CustomerID</u> CustomerName | <u>OrderID</u> <u>ProductID</u> Qty | <u>ProductID</u> ProductName | <u>ProductID</u> ProductName UnitPrice |

Merged
(same primary key)

Orders

| OrderID [PK] | CustomerID | OrderDate | Products Total |
|-----------------|------------|-----------|-------------------|
| A1091 | S009 | 21/7/2011 | 3 |
| A1092 | S010 | 12/1/2011 | 2 |
| A1093 | S010 | 21/8/2011 | 2 |

Derivable
attribute

OrderDetails

| OrderID [PK] | ProductID [PK] | Qty |
|-----------------|-------------------|-----|
| A1091 | S1001 | 100 |
| A1091 | S1003 | 200 |
| A1091 | S1005 | 250 |
| A1092 | S1001 | 10 |
| A1092 | S1004 | 50 |
| A1093 | S1001 | 80 |
| A1093 | S1005 | 90 |

| Orders | Customers | OrderDetails | Products |
|--|-----------------------------------|---|--|
| <u>OrderID</u> CustomerID OrderDate ProductsTotal | <u>CustomerID</u> CustomerName | <u>OrderID</u> <u>ProductID</u> Qty | <u>ProductID</u> ProductName UnitPrice |

Optimization: Remove derivable attributes

| Orders | Customers | OrderDetails | Products |
|---|-----------------------------------|---|--|
| <u>OrderID</u> CustomerID OrderDate ProductsTotal | <u>CustomerID</u> CustomerName | <u>OrderID</u> <u>ProductID</u> Qty | <u>ProductID</u> ProductName UnitPrice |

Removed (derivable Attribute)

Note: Optimization can be done at the end of each normalization step, or after 3NF.

| Orders | Products |
|--|--|
| <u>OrderID</u> CustomerID CustomerName OrderDate <u>ProductID</u> <u>ProductName</u> <u>Qty</u> ProductsTotal | <u>ProductID</u> ProductName UnitPrice |

0NF (top) → 3NF optimized (bottom)

| Orders | Customers | OrderDetails | Products |
|---|-----------------------------------|---|--|
| <u>OrderID</u> CustomerID OrderDate | <u>CustomerID</u> CustomerName | <u>OrderID</u> <u>ProductID</u> Qty | <u>ProductID</u> ProductName UnitPrice |



How “far” should we Normalize?

- For relational databases:
 - 1NF is required, at minimum for practical RDBMS implementations.
 - The majority of the time data models are normalized to 3NF.
 - Sometimes certain tables are left in 1NF or 2NF, for performance or practical reasons.
 - Higher normal forms BCNF, 4NF are rare.
- In General, the Higher the NF of your data model:
 - The more complicated the internal data model
 - The more “programming” required to reproduce the external data model.
 - But, the lesser the chance for data anomalies!!
- It’s a total trade-off: Database complexity vs. data anomalies.

© 2018 National University of Singapore. All Rights Reserved



Agenda


- Introduction to data modelling
- Data analysis
 - Entity relationship diagram (ERD)
 - Attribute analysis
- Data design
 - Normalization
 - **Logical data model**

- Logical data model is a more detailed representation of data
 - Additional Entities (as a result) of normalization
 - Additional Relationships (linking new and existing entities)

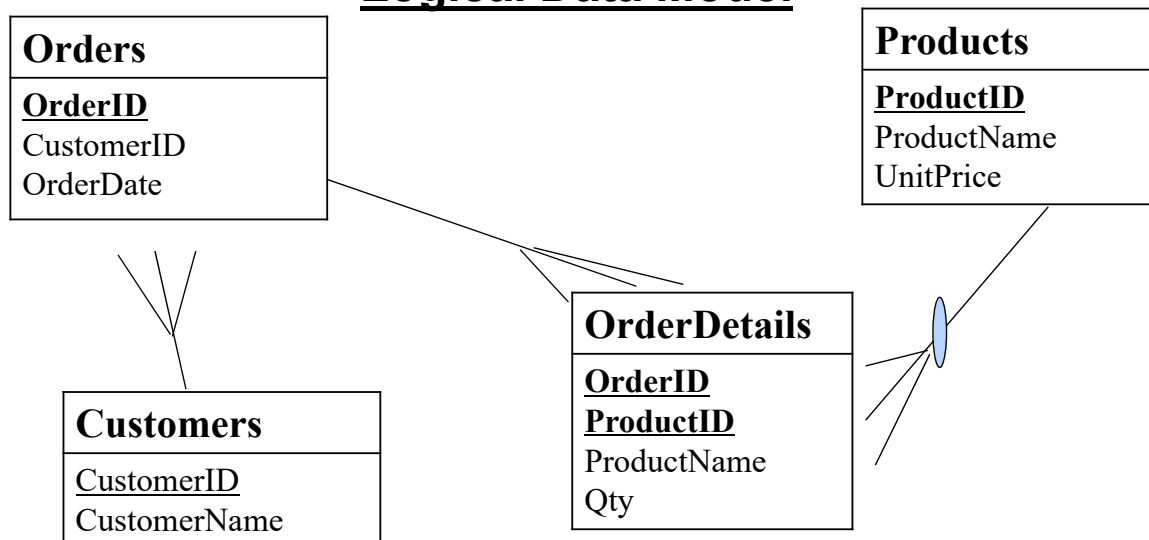


Example

ERD (Before Normalisation)



Logical Data Model



- Design entity relation diagram
- Perform normalization on a conceptual data model
- Design logical data model

Thank you!

Dr TIAN Jing
Email: tianjing@nus.edu.sg