





KE UNIT 3 DATA WAREHOUSING FOR BUSINESS ANALYTICS

DAY 2

Dr TIAN Jing tianjing@nus.edu.sg

© 2018 National University of Singapore. All Rights Reserved















- How can we collect and store large amounts of data?
 - Build tools and data structures to efficiently index and serve data
- How can we efficiently query data?
 - Compile high-level declarative queries into efficient low-level plans
- How can we safely update data?
 - Manage concurrent access to state as it is read and written





- Objective
 - Discuss fundamentals of data management, how to use SQL to query databases for data analytics
 - Not study how to be a database administrator; and how to optimize data query commands
- Outline
 - Relational model
 - SQL query: Data definition language and data manipulation language





- CS145 Introduction to Databases, Stanford University, http://web.stanford.edu/class/cs145/
- Michael Cafarella, Introduction to SQL and the Relational Model, Data Boot Camp, 2014, http://ibugum.github.io/2014-summer-camp/
- Online SQL hands-on exercise, http://sqlfiddle.com/

Page 5



Relational database management systems (RDBMS)





- A schema / database is a collection of logical structures of data or objects
- Some types of objects:
 - Tables
 - stores data / records
 - Indexes
- Mapping from ERD
 - Each Entity -> one relational table
 - Each Attributes-> one column
 - Identifier (key)-> primary key





- All data is stored in tables
 - Organized into rows and columns.
 - Example: employee table.

Emp_No	Emp_Name	IC_No	Dept_No	
179	Chang	P28493	7)
857	Robinson	S95843	4	Rows
342	Bill	T04842	7	

Columns

- All relationship information is presented as data values in tables
 - No ordering

© 2018 National University of Singapore. All Rights Reserved

Page 7





- Each column in a relational database has a datatype
- Common datatypes:

Data Type

CHAR(size) Fixed character data

VARCHAR(size) Variable length character

INT Whole numbers

DECIMAL(p,s) Number with precision and scale

Precision: number of digits in a number.

Scale: number of digits to the right of the decimal point. For example, the number 123.45 has a precision of 5

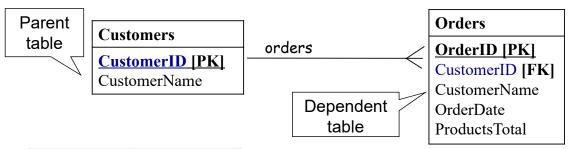
and a scale of 2.

DATE Date field





- A <u>primary key</u> is unique identifiers of the table. At most 1 primary key can be defined on a table.
- A <u>foreign key</u> enforces referential integrity. Foreign key attribute must correspond to an existing primary key value in the parent table (unless the foreign key value is null)



CustomerID	Customer
[PK]	Name
S009	Lynn Wang
S010	Suzan Tan

Order	CustomerID	Order	Products
ID [PK]	[FK]	Date	Total
A1091	S009	21/7/2011	3
A1092	S010	12/1/2011	2



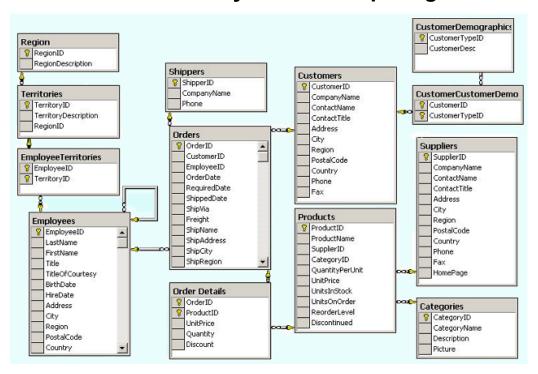


- In RDBMS, databases are illustrated using an ERD (entity relationship diagram)
 - Example of a E-R Model is shown in the next slide.
 - This diagram illustrates Northwind Database provided along with MS SQL Server.
 - This system typically depicts a trading system consisting Sales and Purchases.
 - Each entity is represent by a table : eg. Customers table,
 Employees table, etc. Each row in the table is called the entity data.
 - The lines joining the tables are the relationships.
 - There are a number of relationships between entities in an ERD.
 - 1 to 1 relationship
 - 1 to many relationship
 - Also entity related to itself (Employees)





Northwind Entity Relationship Diagram



© 2018 National University of Singapore. All Rights Reserved





- Interpretation of the Northwind ERD is as follows:
 - Notice that there is an employee id field under the Orders and Employees table.
 - The employee id in Employees table is the primary key while the same id in the Orders table is the foreign key. In other words, for every row in the orders table, that row's employee id must be found in the Employee table.
 - The relationship between the Employee and Orders table is a 1 to many relationship. This means that for each employee id in the Employee table, there could be repetition of the same id in the Orders table.
 - Note that the (One to Many) relationship is determined by the a row of the Employee table with respect to many rows of the Orders table having the same EmployeeID



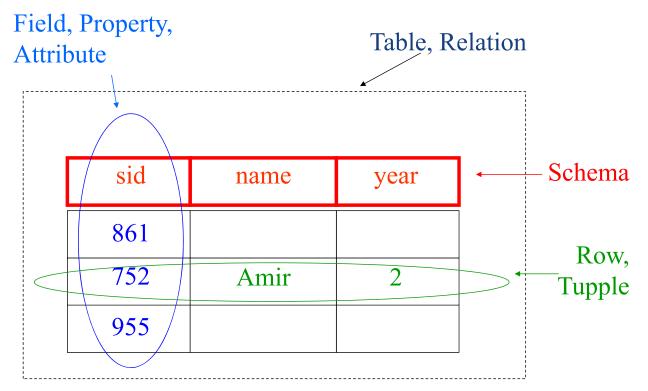


- Interpretation of the Northwind ERD (cont'd):
 - Employees Table has a self referencing relationship. The table has a "report to" column (apart from the employeeID) which requires the employee number to be insert
 - Each employee reports to another employee (I.e., the Boss). Hence each row in the employee table (ie each employee) has a column that requires another EmployeeID for "reporting to" relationship to be established.

employeeid	lastname	firstname	reportsto
1	Davolio	Nancy	2
2	Fuller	Andrew	NULL
3	Leverling	Janet	2
4	Peacock	Margaret	2
5	Buchanan	Steven	2
6	Suyama	Michael	5
7	King	Robert	5
8	Callahan	Laura	2
9	Dodsworth	Anne	5











- Stands for Structured Query Language
- Developed at IBM by Donald D. Chamberlin and Raymond F. Boyce
 - [Chamberlin, Boyce: SEQUEL: A Structured English Query Language. SIGMOD Workshop, Vol. 1 1974: 249-264]
- Originally called SEQUEL
 - Now written SQL but still pronounced "SEQUEL"
- Standardized as ANSI (1986), ISO (1987)

Page 15







- Create the Athlete relation
 - Type constraint enforced when tuples added or modified

CREATE TABLE Athlete (aid INTEGER, name CHAR(30), country CHAR(20), sport CHAR(20));

Create the Olympics relation

CREATE TABLE Olympics (oid INTEGER, year INTEGER, city CHAR(20));

Create the Compete relation

CREATE TABLE Compete (aid INTEGER, oid INTEGER);





Find all athletes from USA:

SELECT *
FROM Athlete;

AID	Name	Country	Sport
1	Mary Lou Retton	USA	Gymnastics
2	Jackie Joyner-Kersee	USA	Track
3	Michael Phelps	USA	Swimming

© 2018 National University of Singapore. All Rights Reserved

Page 17







Can insert a single tuple using

INSERT INTO Athlete (aid, name, country, sport)
VALUES (4, 'Johann Koss', 'Norway', 'Speedskating');





 Can delete all tuples satisfying some condition (e.g., name = Smith)

DELETE
FROM Athlete A
WHERE A.name = 'Smith';

Destroys the relation Olympics.

DROP TABLE Olympics;

© 2018 National University of Singapore. All Rights Reserved

Page 19







Students

name	gpa	age	dept	gender
Sergey Brin	4	40	CS	M
Danah Boyd	4	35	CS	F
Bill Gates	1	60	CS	M
Hillary Mason	4	35	DATASCI	F
Mike Olson	4	50	CS	M
Mark Zuckerberg	4	30	CS	M
Cheryl Sandberg	4	47	BUSINESS	F
Susan Wojcicki	4	46	BUSINESS	F
Marissa Meyer	4	45	BUSINESS	F





- SELECT S.name, S.gpa
 FROM students S
 WHERE S.dept = 'CS'
 [GROUP BY < column list>
 [HAVING < predicate>]]
 [ORDER BY < column list>
- Produce all tuples in the table that satisfy the predicate, output the expressions in the SELECT list.
- Expression can be a column reference, or an arithmetic expression over column references.
- The WHERE clause allows to build arbitrary propositional logic over built-in predicates over attributes
 - Logical operators: AND, OR, NOT
 - Comparisons on numbers/strings (lexicographic): =, !=, >, <, >=, <=,
 - Membership in lists: IN, NOT IN

Page 21







- SELECT S.name, S.gpa
 FROM students S
 WHERE S.dept = 'CS'
 [GROUP BY < column list>
 [HAVING < predicate>]]
 [ORDER BY < column list>
- Try changing the WHERE clause, e.g. 'WHERE S.gpa < 4.0'
- Try changing the SELECT clause, e.g. 'SELECT S.sname, S.gender' (error!)

name	gpa
Sergey Brin	4
Danah Boyd	4
Bill Gates	1
Mike Olson	4
Mark Zuckerberg	4





• SELECT DISTINCT S.name, S.gpa FROM students S
WHERE S.dept = 'CS'
[GROUP BY < column list>
[HAVING < predicate>]]
[ORDER BY < column list>]

• DISTINCT specifies removal of duplicate rows before output

© 2018 National University of Singapore. All Rights Reserved

Page 23





- SELECT S.name, S.gpa, S.age*2 AS a2
 FROM Students S
 WHERE S.dept = 'CS'
 [GROUP BY < column list>
 [HAVING < predicate>]]
 [ORDER BY < column list>]
 - Attributes can be *renamed*
 - Attributes can be invented as *functions of other attributes*

name	gpa	a2
Sergey Brin	4	80
Danah Boyd	4	70
Bill Gates	1	120
Mike Olson	4	100
Mark Zuckerberg	4	60





SELECT S.name, S.gpa, S.age*2 AS a2 FROM Students S WHERE S.dept = 'CS'

ORDER BY S.gpa DESC, S.name ASC, a2 Ascending order by default, but can be overriden

- DESC flag for descending, ASC for ascending
- Can mix and match, lexicographically

name	gpa	a2
Danah Boyd	4	70
Mark Zuckerberg	4	60
Mike Olson	4	100
Sergey Brin	4	80
Bill Gates	1	120

DEMO: http://sqlfiddle.com/#!18/b2a3b/7 © 2018 National University of Singapore. All Rights Reserved

Page 25





SELECT [DISTINCT] AVG(S.gpa) FROM Students S WHERE S.dept = 'CS'[ORDER BY < column list>];

Result is 3.4

Common aggregate functions

SUM(C) – sum over all numbers in C COUNT(C) - number of rows in C AVG(C) - SUM(C)/COUNT(C)MAX(C) – largest value MIN(C) – smallest value





- SELECT [DISTINCT] AVG(S.gpa), S.dept
 FROM Students S
 [WHERE < predicate >]
 GROUP BY S.dept
 [HAVING < predicate >]
 [ORDER BY < column list >];
- Partition table into groups with same GROUP BY column values
- Produce an aggregate result per group
- Note: can put grouping columns in SELECT list

	dept
4	BUSINESS
3.4	CS
4	DATASCI

DEMO: http://sqlfiddle.com/#!18/b2a3b/16

© 2018 National University of Singapore. All Rights Reserved

Page 27

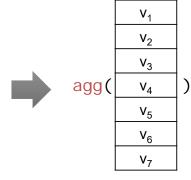
Grouping idea





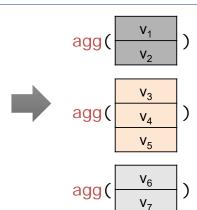
SELECT agg(B)

А	В
а	v ₁
а	V_2
b	V ₃
b	V ₄
b	V ₅
С	v ₆
С	V ₇
-	



SELECT agg(B)
GROUP BY A

Α	В
а	V ₁
а	V_2
b	V ₃
b	V_4
b	V ₅
С	v ₆
С	V ₇







- SELECT [DISTINCT] AVG(S.gpa), S.dept FROM Students S
 [WHERE < predicate >]
 GROUP BY S.dept HAVING COUNT(*) > 2
 [ORDER BY < column list >];
- The HAVING predicate filters groups
- HAVING is applied *after* grouping and aggregation
 - Hence can contain anything that could go in the SELECT list, i.e. aggs or GROUP BY columns
- HAVING can only be used in aggregate queries
- It's an optional clause

	dept
4	BUSINESS
3.4	CS

DEMO: http://sqlfiddle.com/#!18/b2a3b/17 © 2018 National University of Singapore. All Rights Reserved

Page 29

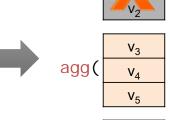
Group selection





SELECT agg(B)
GROUP BY A
HAVING A>a

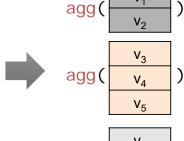
Α	В
а	v ₁
а	V ₂
р	V_3
b	V ₄
р	V ₅
С	V ₆
С	V ₇

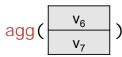


000(V_6	\
agg(V ₇)

SELECT agg(B)
GROUP BY A

Α	В	
а	V ₁	
а	V ₂	
b	V ₃	
b	V_4	
b	V ₅	
C	V ₆	
С	V ₇	









• SELECT S.dept, AVG(S.gpa), COUNT(*) FROM Students S

WHERE S.gender = 'F'
GROUP BY S.dept
HAVING COUNT(*) >= 2
ORDER BY S.dept;

dept		
BUSINESS	4	3

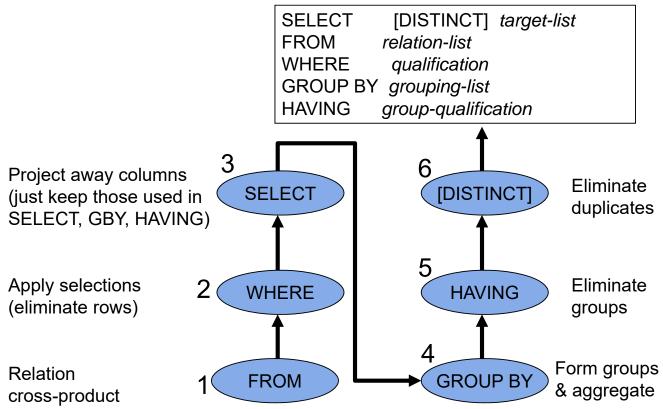
name	gpa	age	dept	gender
Sergey Brin	4	40	CS	M
Danah Boyd	4	35	CS	F
Bill Gates	1	60	CS	M
Hillary Mason	4	35	DATASCI	F
Mike Olson	4	50	CS	M
Mark Zuckerberg	4	30	CS	M
Cheryl Sandberg	4	47	BUSINESS	F
Susan Wojcicki	4	46	BUSINESS	F
Marissa Meyer	4	45	BUSINESS	F

DEMO: http://sqlfiddle.com/#!18/b2a3b/18

© 2018 National University of Singapore. All Rights Reserved











- SELECT [DISTINCT] < column expression list>
 FROM < table1 [AS t1], ..., tableN [AS tn]>
 [WHERE < predicate>]
 [GROUP BY < column list>
 [HAVING < predicate>]]
 [ORDER BY < column list>];
 - 1. FROM: compute *cross product* of tables.
 - 2. WHERE: Check conditions, discard tuples that fail.
 - 3. SELECT: Specify desired fields in output.
 - 4. DISTINCT (optional): eliminate duplicate rows.

Page 33





All pairs of tuples, concatenated

IJ

uid	uname	followers	age
1	Anurag	2224	22
2	Ryan	12	39
3	Rolando	2	27
4	Valerie	5034	19

P

uid	cid	time
1	102	2018-01-16 09:00
2	102	2018-01-17 09:02
1	101	2018-01-22 23:00

U.uid	U.uname	U.followers	U.age	P.uid	P.cid	P.time
1	Anurag	2224	22	1	102	2018-01-16 09:00
1	Anurag	2224	22	2	102	2018-01-17 09:02
1	Anurag	2224	22	1	101	2018-01-22 23:00
2	Ryan	12	39	1	102	2018-01-16 09:00
				•••		





uid	uname	followers	age
1	Anurag	2224	22
2	Ryan	12	39
3	Rolando	2	27
4	Valerie	5034	19

Users

cid	cname	bkcolor
101	homeworks	red
102	midterms	magenta
103	lectures	red

Channels

Posts

uid	cid	time
1	101	2018-01-22 23:00:00
1	102	2018-01-16 09:00:00
2	102	2018-01-17 09:02:00

© 2018 National University of Singapore. All Rights Reserved

Page 35





SELECT Users.uid, uname FROM Users, Posts WHERE Users.uid = Posts.uid

uid	uname
1	Anurag
1	Anurag
2	Ryan

SELECT U.uid, uname FROM Users AS U, Posts AS P WHERE U.uid=P.uid

DEMO: http://sqlfiddle.com/#!18/ed664/1





- Nesting: one query is nested in another query as a relation/value component
- The nested query is called a subquery
- Where are we nesting?
 - SELECT
 - Select a value from a subquery
 - FROM
 - Use a subquery instead of an existing relation
 - WHERE
 - Conditions phrased via subqueries

Page 37







Names of users who've posted on channel #102:

subquery

SELECT U.uname
FROM Users U
WHERE U.uid IN
(SELECT P.uid
FROM Posts P
WHERE P.cid=102);

uname	
Anurag	
Ryan	

SELECT P.uid FROM Posts P WHERE P.cid=102;

SELECT U.uname FROM Users U WHERE U.uid IN ('1', '2');

uid	uname
1	Anurag
2	Ryan

DEMO: http://sqlfiddle.com/#!18/a7ca4/7





Names of users who've **not** posted on channel #103:

SELECT U.uname
FROM Users U
WHERE U.uid NOT IN
(SELECT P.uid
FROM Posts P
WHERE P.cid=103)

uname	
Anurag	
Ryan	
Rolando	
Valerie	

DEMO: http://sqlfiddle.com/#!18/a7ca4/9

© 2018 National University of Singapore. All Rights Reserved





- INNER is default
- Inner join is akin to what we've learned so far, just with different syntax.

```
SELECT (column_list)
FROM table_name
[INNER | {LEFT |RIGHT | FULL } {OUTER}] JOIN table_name
ON qualification_list
WHERE ...
```





Both are equivalent!

SELECT u.uid, u.uname, p.uid

WHERE u.uid = p.uid

FROM Users u, Posts p

AND u.age > 20;

SELECT u.uid, u.uname, p.uid

FROM Users u **INNER JOIN** Posts p

ON u.uid = p.uid

AND u.age > 20;

uid	uname	uid
1	Anurag	1
1	Anurag	1
2	Ryan	2

DEMO: http://sqlfiddle.com/#!18/a7ca4/11

© 2018 National University of Singapore. All Rights Reserved

Page 41



🛊 Tables used in example





Sailors

sid	sname	rating	age
22	Dustin	7	45
31	Lubber	8	55.5
95	Bob	3	63.5

Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue

Reserves

sid	bid	day
22	101	10/10/96
95	103	11/12/96





Returns all matched rows, <u>and preserves</u> <u>all unmatched rows from the table on the left</u> of the join clause (use nulls in fields of non-matching tuples)

sid	sname	bid
22	Dustin	101
95	Bob	103
31	Lubber	(null)

SELECT s.sid, s.sname, r.bid FROM Sailors s LEFT OUTER JOIN Reserves r ON s.sid = r.sid;

DEMO: http://sqlfiddle.com/#!17/78155/2

© 2018 National University of Singapore. All Rights Reserved

Page 43







Returns all matched rows, <u>and preserves</u> <u>all unmatched rows from the table on the right</u> of the join clause (use nulls in fields of non-matching tuples)

sid	sname	bid
22	Dustin	101
95	Bob	103
31	Lubber	(null)

SELECT s.sid, s.sname, r.bid FROM Reserves r RIGHT OUTER JOIN Sailors s ON r.sid = s.sid;

DEMO: http://sqlfiddle.com/#!17/78155/4





Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

sid	bid	bname
22	101	Nina
95	103	Santa Maria
(null)	102	Pinta

SELECT r.sid, b.bid, b.bname FROM Reserves r FULL OUTER JOIN Boats b ON r.bid = b.bid

DEMO: http://sqlfiddle.com/#!17/78155/5

© 2018 National University of Singapore. All Rights Reserved





- Field values are sometimes unknown
 - SQL provides a special value NULL for such situations.
 - Every data type can be NULL
- The presence of null complicates many issues. E.g.:
 - Selection predicates (WHERE)
 - Aggregation
- NULLs also come from various joins





Sailors

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19
5	SpongeBob	11	2
11	Jack Sparrow	(null)	35

Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue

Reserves

sid	bid	day
1	101	10/01/2017
1	102	9/12/2017
2	102	9/13/2017

© 2018 National University of Singapore. All Rights Reserved

Page 47

NULL in WHERE clause



WHERE clause eliminates rows that don't evaluate to true

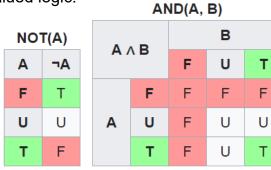
SELECT * FROM sailors WHERE rating > 8;

SELECT * FROM sailors WHERE rating > 8 AND 'True'='True';

SELECT * FROM sailors WHERE rating > 8 OR 'True'='True';

SELECT * FROM sailors WHERE NOT (rating > 8);

Three-valued logic:



OR(A, B)



DEMO: http://sqlfiddle.com/#!18/0f6bd/21 © 2018 National University of Singapore. All Rights Reserved







NULL in WHERE clause

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
5	SpongeBob	11	2
sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
5	SpongeBob	11	2
sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19
5	SpongeBob	11	2
11	Jack Sparrow	(null)	35
sid	sname	rating	age
3	Garfield	1	27
4	Bob	5	19

Page 49





NULL values generally ignored when computing aggregates

6

5

38

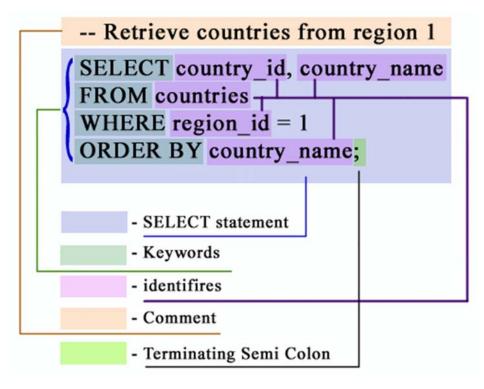
7

SELECT count(*) FROM sailors; SELECT count(rating) FROM sailors; SELECT sum(rating) FROM sailors; SELECT avg(rating) FROM sailors;

sid	sname	rating	age	
1	Popeye	10	22	
2	OliveOyl	11	39	
3	Garfield	1	27	
4	Bob	5	19	
5	SpongeBob	11	2	
11	Jack Sparrow	(null)	35	







Source: https://www.w3resource.com/sql/sql-syntax.php

© 2018 National University of Singapore. All Rights Reserved

Page 51



掉 Summary: SQL command





SOL SELECT STATEMENTS SELECT * FROM tbl Select all rows and columns from table tbl SELECT c1,c2 FROM tbl Select column c1, c2 and all rows from table tbl SELECT c1,c2 FROM tbl WHERE conditions ORDER BY c1 ASC, c2 DESC Select columns c1, c2 with where conditions and from table tbl order result by column c1 in ascending order and c2 in descending order

SELECT DISTINCT c1, c2

FROM thi

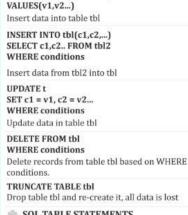
Select distinct rows by columns c1 and c2 from

SELECT c1, aggregate(expr) FROM tbl **GROUP BY c1**

Select column c1 and use aggregate function on expression expr, group columns by column c1.

SELECT c1, aggregate(expr) AS c2 FROM tbl CROUP BY c1

HAVING c2 > v Select column c1 and c2 as column alias of the result of aggregate function on expr. Filter group of records with c2 greater than value v



SQL TABLE STATEMENTS

SQL UPDATE TABLE

INSERT INTO tbl(c1,c2,...)

CREATE TABLE tbl(c1 datatype(length) c2 datatype(length)

PRIMARY KEY(c1)

Create table tbl with primary key is c1

DROP TABLE thi

Remove table tbl from database.

ALTER TABLE thi

ADD COLUMN c1 datatype(length)

Add column c1 to table tbl

ALTER TABLE thi DROP COLUMN c1

Drop column c1 from table tbl

SQL JOIN STATEMENTS

SELECT * FROM tbl1

INNER JOIN tbl2 ON join-conditions

Inner join table tbl1 with tbl2 based on joinconditions.

SELECT * FROM tbl1

LEFT JOIN tbl2 ON join-conditions

Left join table tbl1 with tbl2 based on join-

SELECT * FROM tbl1

RIGHT JOIN tbl2 ON join-conditions

Right join table tbl1 with tbl2 based on joinconditions.

SELECT * FROM tbl1

RIGHT JOIN tbl2 ON join-conditions

Full outer join table tbl1 with tbl2 based on joinconditions.





- Write SQL queries to find
 - 1. Team names for all teams with attendance more than 2,000,000
 - 2. Player ID and home stadium for all Allstars
 - 3. TeamID, attendance values for teams that had an all-star player ORDERED BY ATTENDANCE

SELECT name FROM Teams WHERE attendance > 2000000;

SELECT playerID, park FROM Allstars, Teams WHERE Allstars.teamID = Teams.teamID;

SELECT DISTINCT Allstars.teamID, attendance FROM Teams, Allstars WHERE Teams.teamID = Allstars.teamID ORDER BY attendance DESC;

DEMO: http://sqlfiddle.com/#!18/66363/111

© 2018 National University of Singapore. All Rights Reserved

Page 53





- Write SQL queries to find
 - 1. Average attendance for all teams
 - 2. Average attendance among teams that had an allstar player

SELECT AVG(attendance) FROM Teams;

SELECT AVG(attendance) FROM Teams, Allstars WHERE Teams.teamID = Allstars.teamID;

DEMO: http://sqlfiddle.com/#!18/66363/114





- 1. Show all teamlds that had an all-star, along with number of all-star players;
- 2. Show all team names that had an all-star, along with number of all-star players;
- Show all team names that had an all-star, along with number of all-star players, SORTED IN DESCENDING ORDER OF NUM ALLSTARS AND: only show teams with at least 2 players

SELECT teamID, COUNT(*) FROM Allstars GROUP BY teamID;

SELECT name, COUNT(Allstars.playerID) FROM Allstars, Teams WHERE Allstars.teamID = Teams.teamID GROUP BY name;

SELECT name, COUNT(Allstars.playerID) AS playerCount FROM Allstars, Teams WHERE Allstars.teamID = Teams.teamID GROUP BY name HAVING COUNT(Allstars.playerID) >= 2 ORDER BY playerCount DESC;

DEMO: http://sqlfiddle.com/#!18/66363/113

© 2018 National University of Singapore. All Rights Reserved

Page 55





Thank you!

Dr TIAN Jing
Email: tianjing@nus.edu.sg