



COLUMN FAMILY DATABASE

CASSANDRA

Yunghans Irawan (yirawan@nus.edu.sg)

- Column Family Database
- Cassandra
- Terminology
- Data Types
- Data Modeling
- Architecture
- Summary

- Store data in column families
 - Many column associated with a row key
- Column families are groups of related data that is often accessed together
- Arguably more similar to RDBMS compared to other types of NoSQL databases



Column Family Databases

Name	Initial Release	Latest Version	License
HBase	2008	1.4.3 April 2018	Open Source
Cassandra	2008	3.11.2 Feb 2018	Open Source



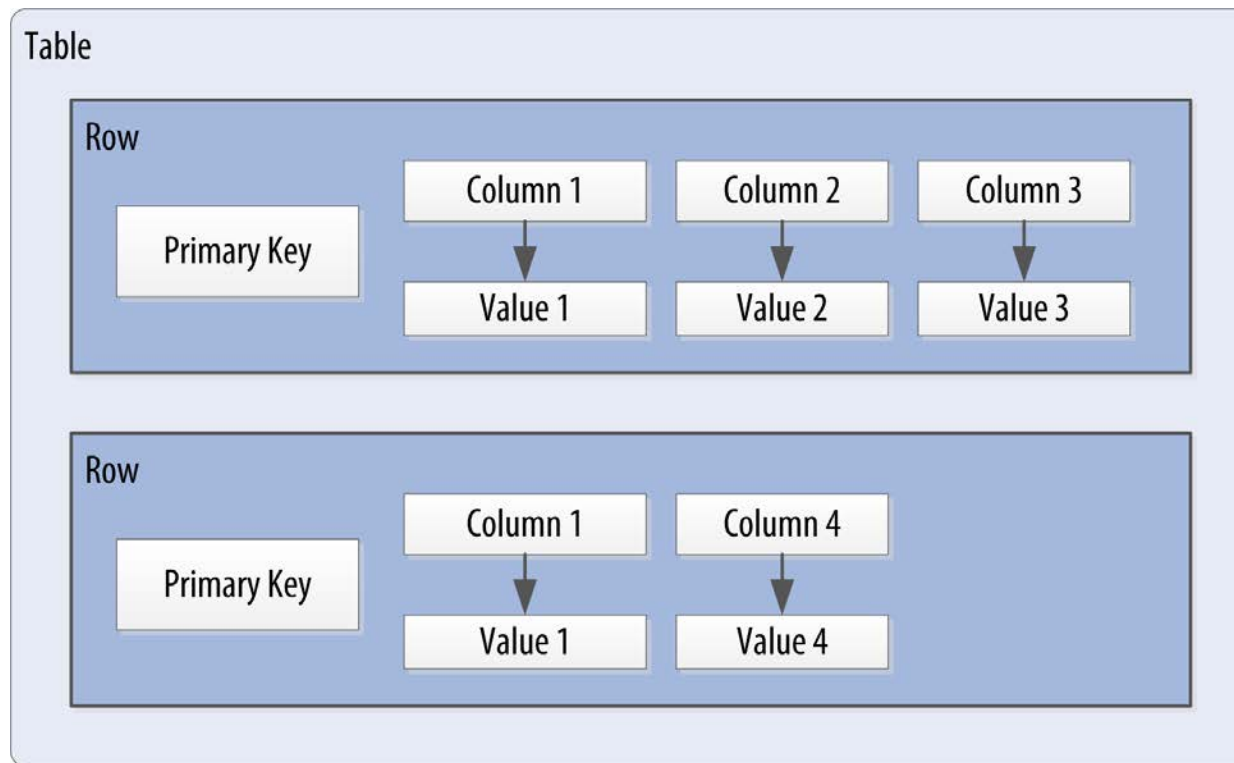
Cassandra

- Open source Apache Project
- Originated at Facebook in 2007 to solve its inbox search problem
 - Huge volumes of data
 - Many random reads
 - Many simultaneous random writes
- Paper on Cassandra is published in 2009
- Project is led and supported by DataStax
 - Provide enterprise versions, integrations with other technologies and product support

- Can support high rate of write
- Scalable write and read throughput, suitable for large deployment
- Similarity with RDBMS concepts with richer data types
- Many configurable parameters depending to customize the behavior based on needs
 - E.g. replication and consistency requirements
- Can support very large number of “columns”
 - 2 billions variable columns is the limit
 - Quote mark because the behavior is kind of different from traditional RDBMS column we are familiar with

- Weaker ACID transactions property
 - Starts with no transaction support
 - More scenarios are being supported as the product matures
- Requires more design effort
 - Similar to RDBMS but not quite the same
- No joins
- No referential integrity enforcement

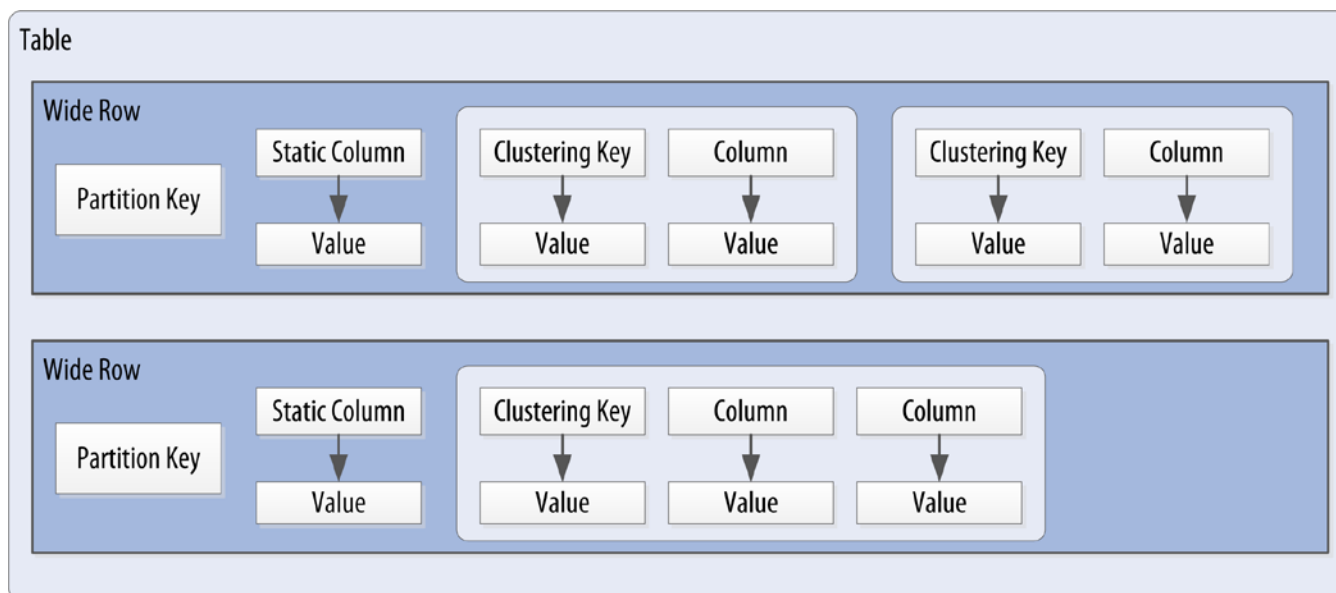
- A logical collection of similar rows





Compound Key

- Compound key is made of:
 - 1 partition key
 - 1 or more clustering key
- Partition key determines which node stores the data





Static column

- Can be quite a complicated concept looking from RDBMS perspective
- Behave like 'static' keyword in C# or Java
 - If you're from developer background
- It means that there's only one of such cell for one partition key, regardless of the number of clustering key
 - Refer to the previous diagram
- Example – Attendance taking
 - Compound key: ClassID (P) + Date (C)+ Student ID (C)
 - Static column: Instructor
 - Assuming instructor depends only on the class and not on the date/student

- Column
 - Name and value pair
- Row
 - Container for columns referenced by a primary key
- Table
 - Container for rows
- Keyspace
 - Container for tables
- Cluster
 - Container for keyspaces that span one or more nodes



Comparison of terms

RDBMS	Cassandra
Cluster	Cluster
Database	Keyspace
Table	Table
Row	Row
Column	Column



Cassandra Data Types

CQL Type	Constants	Description
ascii	strings	US-ASCII character string
bigint	integers	64-bit signed long
blob	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
boolean	booleans	true or false
counter	integers	Distributed counter value (64-bit long)
date	strings	Date string, such as 2015-05-03



Cassandra Data Types

CQL Type	Constants	Description
decimal	integers, floats	Variable-precision decimal Java type Note: When dealing with currency, it is a best practice to have a currency class that serializes to and from an int or use the Decimal form.
double	integers, floats	64-bit IEEE-754 floating point Java type
float	integers, floats	32-bit IEEE-754 floating point Java type



Cassandra Data Types

CQL Type	Constants	Description
frozen	user-defined types, collections, tuples	A frozen value serializes multiple components into a single value. Non-frozen types allow updates to individual fields. Cassandra treats the value of a frozen type as a blob. The entire value must be overwritten. Note: Cassandra no longer requires the use of frozen for tuples :frozen <tuple < int , tuple<text, double >>>
inet	strings	IP address string in IPv4 or IPv6 format, used by the python-cql driver and CQL native protocols
int	integers	32-bit signed integer
list	n/a	A collection of one or more ordered elements: [literal, literal, literal]



Cassandra Data Types

CQL Type	Constants	Description
list	n/a	A collection of one or more ordered elements: [literal, literal, literal]
map	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
set	n/a	A collection of one or more elements: { literal, literal, literal }
smallint	integers	2 byte integer
text	strings	UTF-8 encoded string
time	strings	Time string, such as 13:30:54.234
timestamp	integers, strings	Date plus time, encoded as 8 bytes since epoch
timeuuid	uuids	Version 1 UUID only



Cassandra Data Types

CQL Type	Constants	Description
tinyint	integers	1 byte integer
tuple	n/a	Cassandra 2.1 and later. A group of 2-3 fields.
uuid	uuids	A UUID in standard UUID format
varchar	strings	UTF-8 encoded string
varint	integers	Arbitrary-precision integerJava type



Cassandra compared to RDBMS

- Denormalization is normal
 - Logical side effect of the first two characteristic
- Cannot use column that is not indexed as part of the query criteria
- Query-first design



Secondary Index

- Allow query to use the column as criteria
- Can be applied to simple type column as well as collection columns

```
> SELECT * FROM user WHERE last_name = 'Nguyen';
InvalidRequest: code=2200 [Invalid query] message="No supported
  secondary index found for the non primary key columns restrictions"

> CREATE INDEX ON user ( last_name );

> SELECT * FROM user WHERE last_name = 'Nguyen';
first_name | last_name
-----+-----
      Bill |   Nguyen
(1 rows)

> CREATE INDEX ON user ( emails );
> CREATE INDEX ON user ( phone_numbers );
```

Data Modeling for Cassandra

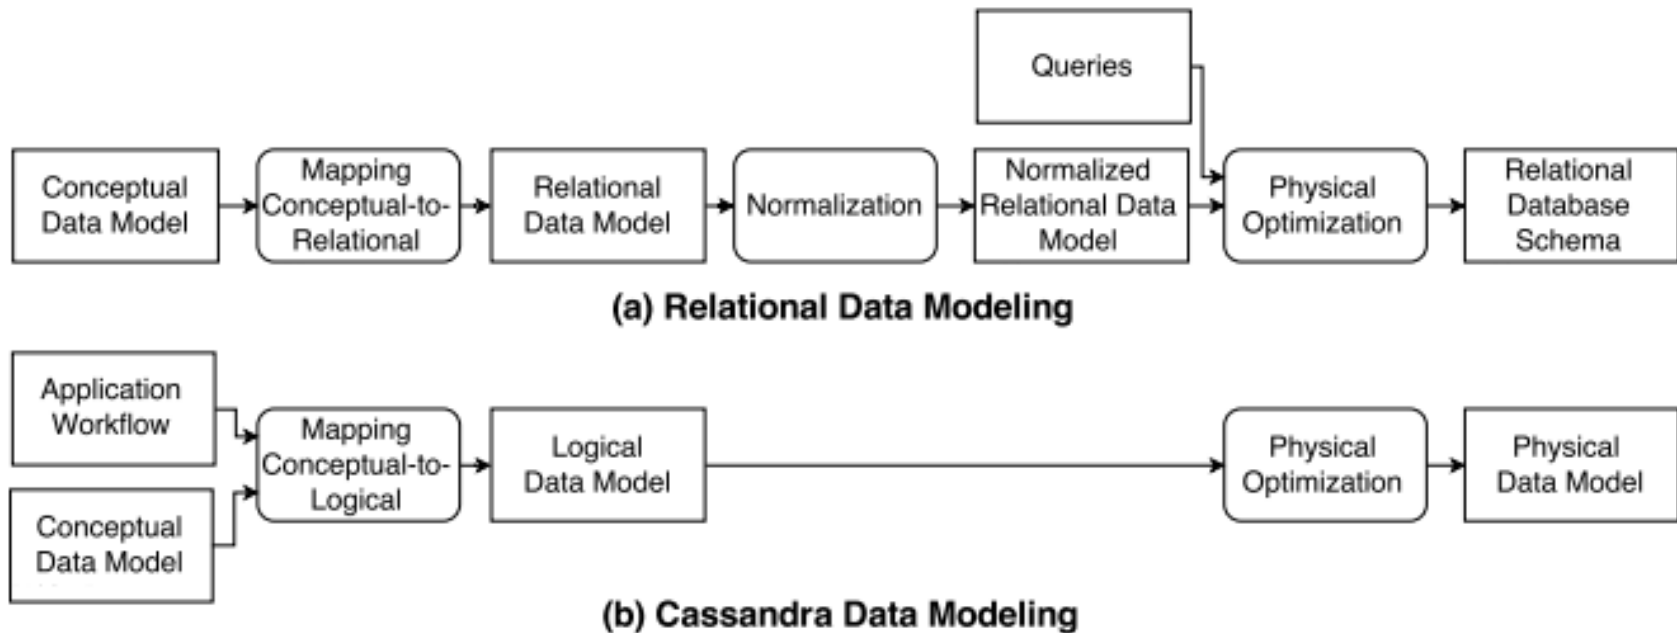
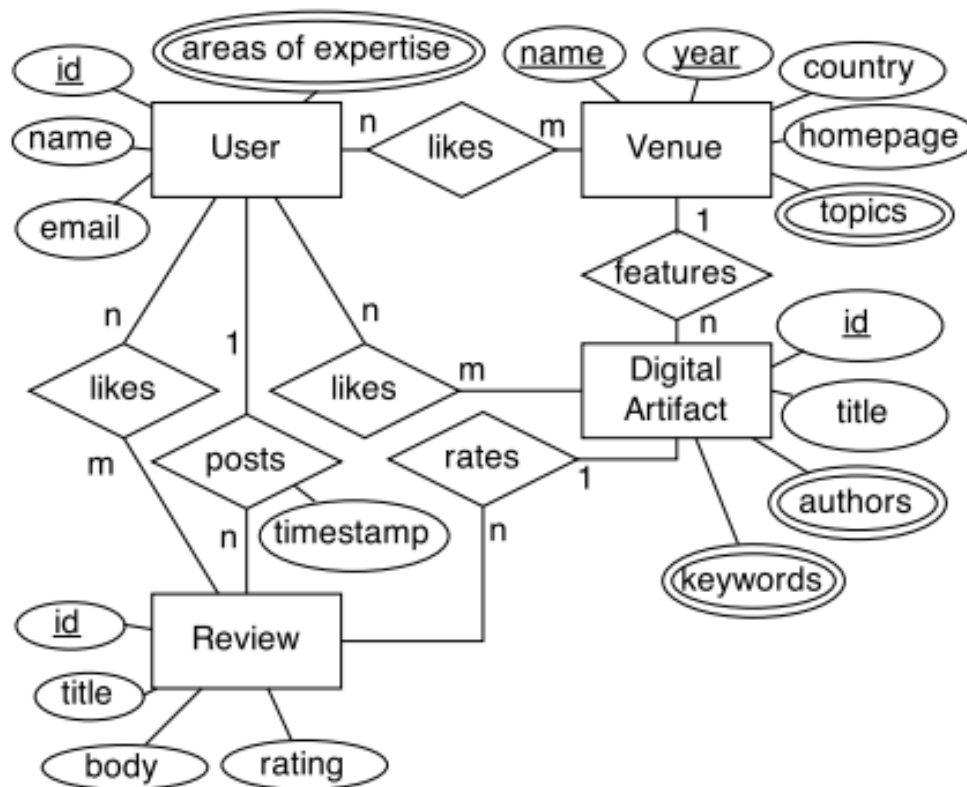


Fig. 1: Traditional data modeling compared with our proposed methodology for Cassandra.



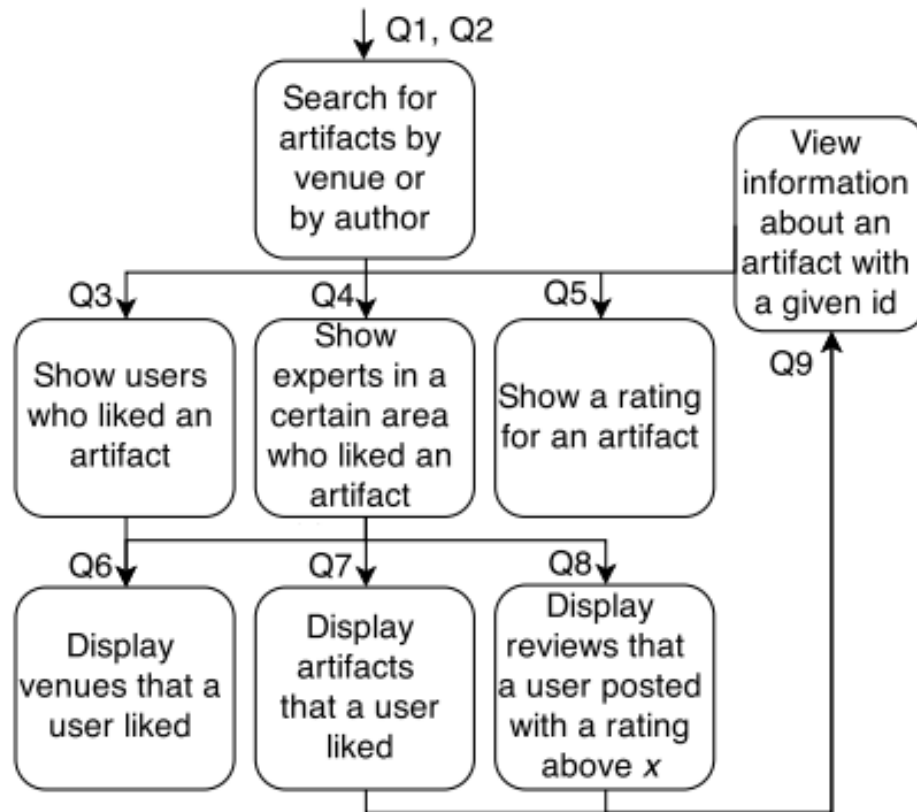
Conceptual Data Model



(a) Conceptual data model



Application Workflow



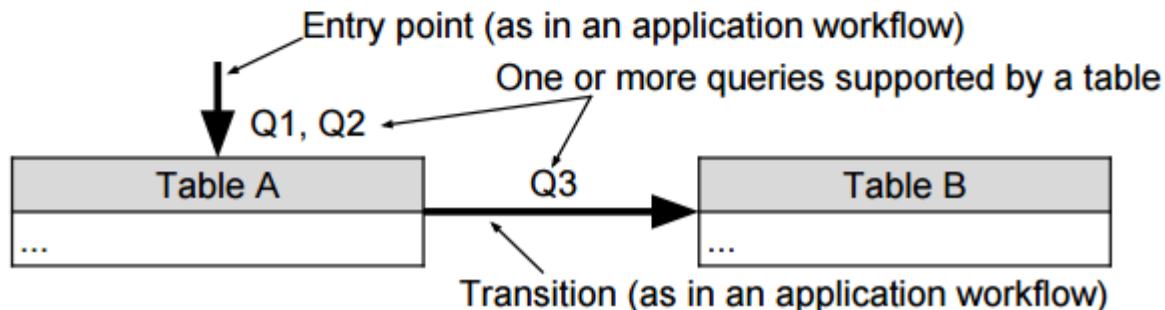
- Q1: Find artifacts published in a venue with a given name after a given year. Order results by year (DESC).
- Q2: Find artifacts published by a given author. Order results by year (DESC).
- Q3: Find users who liked a given artifact.
- Q4: Find users who liked a given artifact and who have expertise in a certain area.
- Q5: Find an average rating of a given artifact.
- Q6: Find venues that a given user liked.
- Q7: Find artifacts published after a certain year that a given user liked. Order results by year (DESC).
- Q8: Find reviews posted by a given user with a rating $\geq x$. Order results by rating (DESC).
- Q9: Find information about an artifact with a given id.

(b) Application workflow

Chebotko Diagram Notation

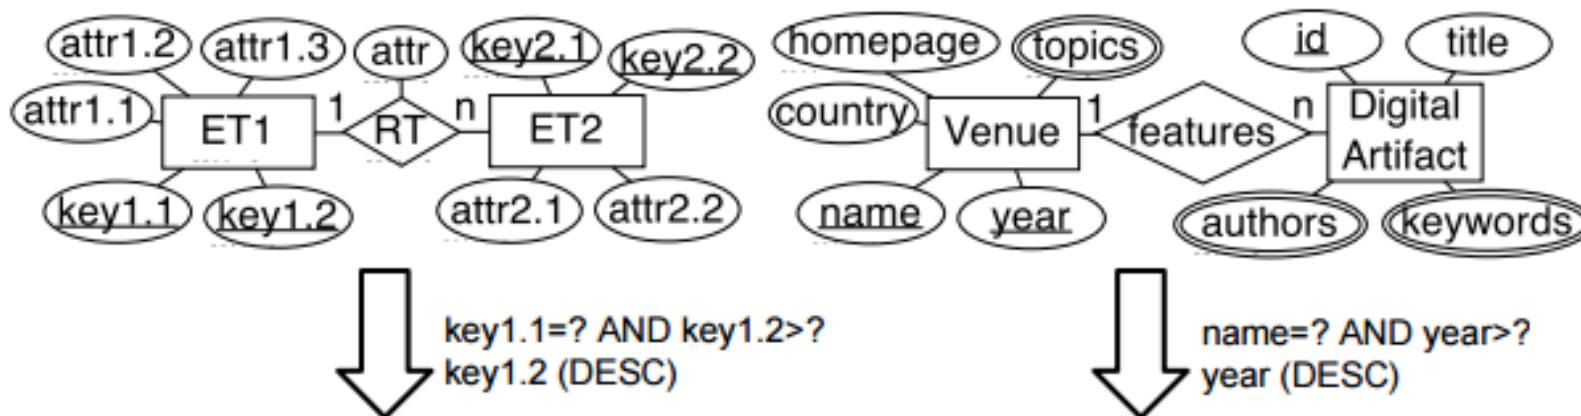
Table schema

Table Name			
column name 1	CQL-Type	K	← Partition key column
column name 2	CQL-Type	C↑	← Clustering key column (ASC)
column name 3	CQL-Type	C↓	← Clustering key column (DESC)
column name 4	CQL-Type	S	← Static column
column name 5	CQL-Type	IDX	← Secondary index column
column name 6	CQL-Type	++	← Counter column
[column name 7]	CQL-Type		← Collection column (list)
{column name 8}	CQL-Type		← Collection column (set)
<column name 9>	CQL-Type		← Collection column (map)
column name 10	CQL-Type		← Regular column





Mapping into Logical



ET2_by_ET1	
key1.1	K
key1.2	C↓
key2.1	C↑
key2.2	C↑
attr2.1	
attr2.2	
attr	

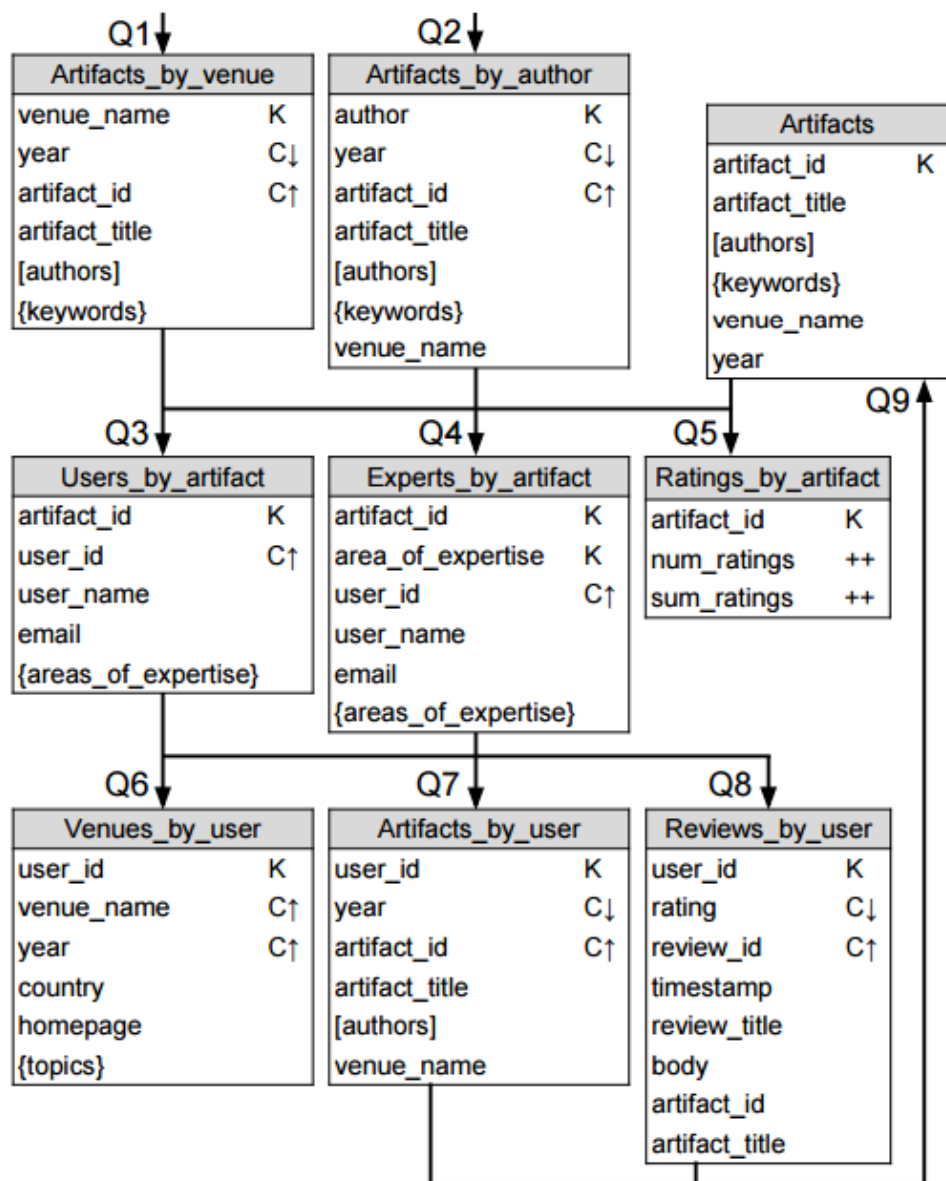
(a) Sample mapping pattern

Artifacts_by_venue	
venue_name	K
year	C↓
artifact_id	C↑
artifact_title	
[authors]	
{keywords}	

(b) Example mapping pattern application

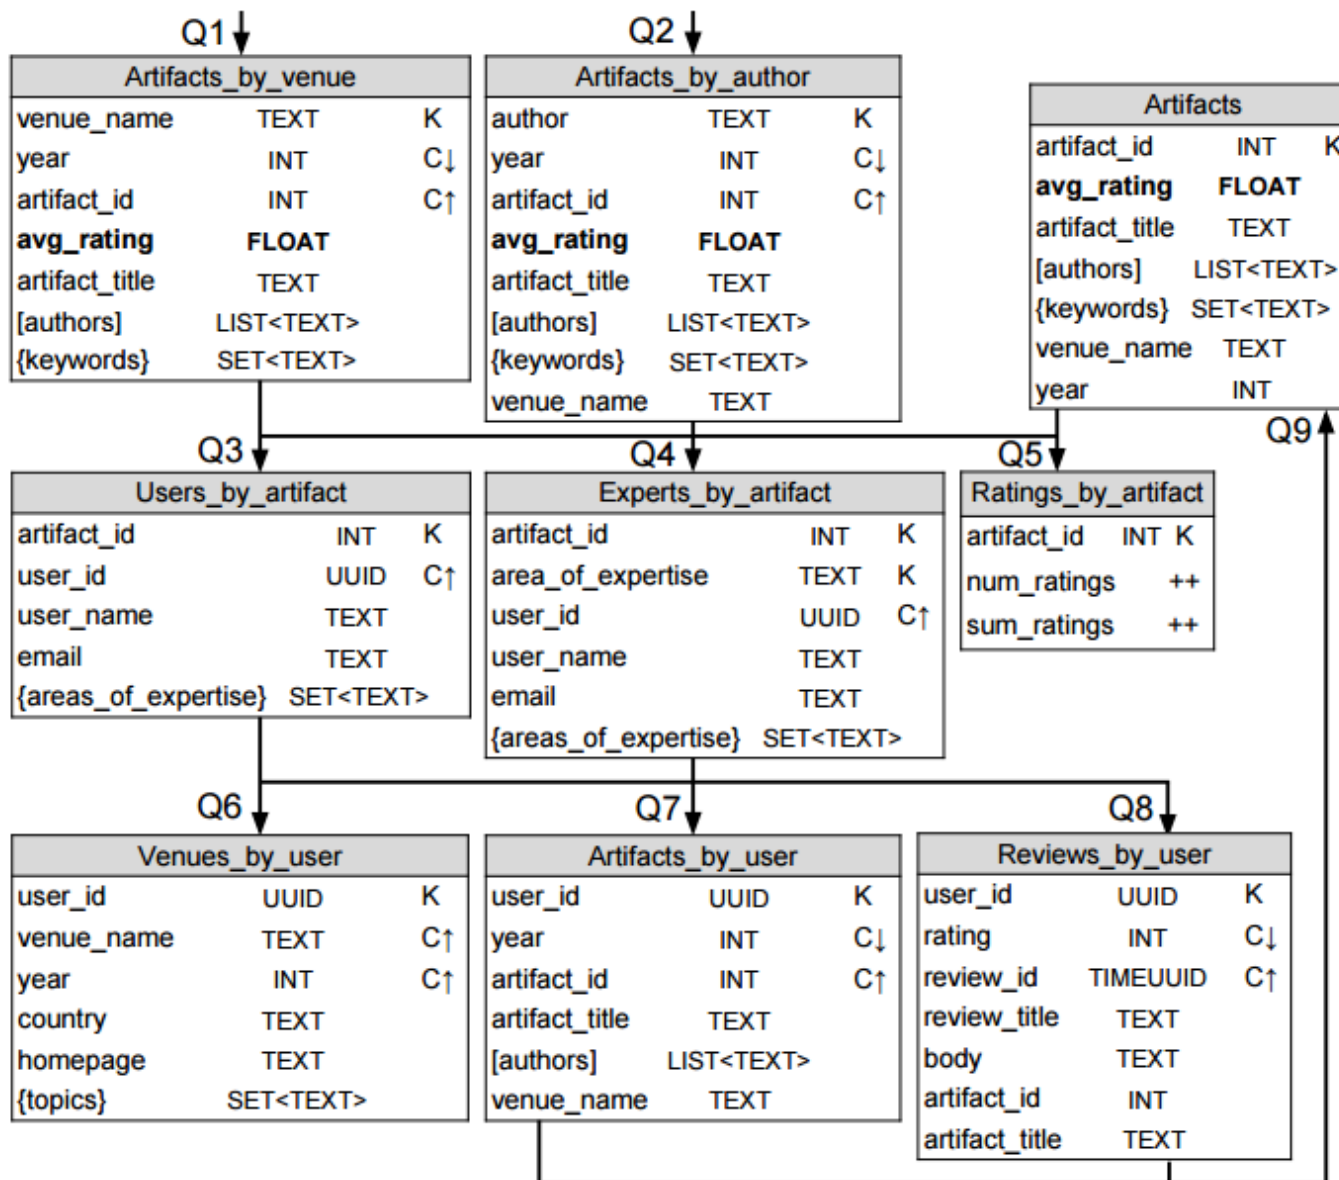


Logical Diagram









Physical Diagram





Cassandra Architecture

- Ring based replication
- No single point of failure
- Key hashes to a location in the ring
- Virtual nodes
 - Allow non-contiguous assignment of ranges
 - Data is split into more ranges
- Cluster rebalancing is automatically done when adding or removing nodes





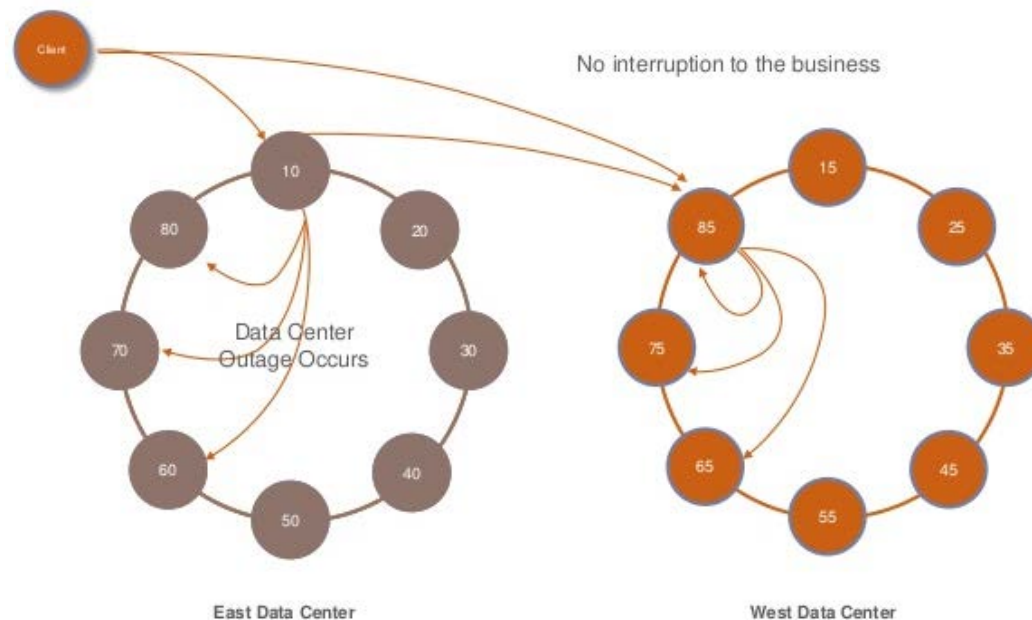
Consistency Level

- Consistency level for read and write is tunable
- Write consistency levels
 - ALL, EACH_QUORUM, QUORUM, LOCAL_QUORUM, ONE, TWO, THREE, LOCAL_ONE, ANY
- Read consistency levels
 - ALL, EACH_QUORUM, QUORUM, LOCAL_QUORUM, ONE, TWO, THREE, LOCAL_ONE, SERIAL, LOCAL_SERIAL
- <https://docs.datastax.com/en/cassandra/3.0/cassandra/dml/dmlConfigConsistency.html>



Data Center Aware Replication

- Instances can be tagged with the datacenter ID to create multi data center deployment
- Consistency levels are data center aware




Source: <http://image.slidesharecdn.com/nycjavameetupdsepresentation-140701094225-phpapp01/95/datatax-nyc-java-meetup-cassandra-with-java-8-638.jpg?cb=1404207863>




Example: Read with QUORUM

Single data center cluster with 3 replica nodes and consistency set to QUORUM




Multiple data center cluster with 3 replica nodes and consistency set to LOCAL_QUORUM



Source: https://teddyma.gitbooks.io/learncassandra/content/client/read_requests.html



Example: Write



- Column-family database offers performance and scalability beyond what traditional RDBMS offers
- The trade-off is extra complexity at design and application development
- Sophisticated replication technique
 - Data center aware
 - Seamless partition rebalancing



References

- Cassandra Documentation
 - <http://docs.datastax.com/>
- Original Cassandra Paper
 - <http://www.cs.cornell.edu/projects/ladis2009/papers/lakshman-ladis2009.pdf>
- Cassandra Reference Card
 - http://www.datastax.com/wp-content/uploads/2013/03/cql_3_ref_card.pdf
- Cassandra: The Definitive Guide, 2nd Edition
 - Jeff Carpenter, Eben Hewitt, O'Reilly Media, 2016
- A Big Data Modeling Methodology for Apache Cassandra
 - <http://www.cs.wayne.edu/andrey/papers/TR-BIGDATA-05-2015-CKL.pdf>