

Workshop - KE

Time Series Forecasting (AR & MA) using R

Data Source:

1. Hotel Data - R.csv and t6-4 hotel_subset.csv (same data with and without dummy)
2. data_1995.csv
3. AmtrakBig.csv
4. cola-data.csv
5. Electricity raw data.csv

In [4]:

```
# setup the infrastructure
pacman::p_load(tidyverse, lubridate, zoo, forecast, fUnitRoots)
```

In [5]:

```
# set path to data folder on your computer
# pay attention to the '/'
setwd("C:/Users/isspcs/Desktop/workshop-data")
```

A. Linear Regression

Let us use the hotel data we saw during lecture to quickly develop a linear regression model

In [78]:

```
data = read.csv('Hotel Data - R.csv')
```

In [79]:

```
# check the data types of all columns ?  
sapply(data, class)
```

```
t  
'integer'  
Yt  
'integer'  
Yt_trans  
'numeric'  
t.1  
'integer'  
S1  
'integer'  
S2  
'integer'  
S3  
'integer'  
S4  
'integer'  
S5  
'integer'  
S6  
'integer'  
S7  
'integer'  
S8  
'integer'  
S9  
'integer'  
S10  
'integer'  
S11  
'integer'
```

In [80]:

```
# get strucutre of data ?  
str(data)
```

```
'data.frame':  168 obs. of  15 variables:  
 $ t      : int  1 2 3 4 5 6 7 8 9 10 ...  
 $ Yt      : int  501 488 504 578 545 632 728 725 585 542 ...  
 $ Yt_trans: num  4.73 4.7 4.74 4.9 4.83 ...  
 $ t.1     : int  1 2 3 4 5 6 7 8 9 10 ...  
 $ S1      : int  1 0 0 0 0 0 0 0 0 0 ...  
 $ S2      : int  0 1 0 0 0 0 0 0 0 0 ...  
 $ S3      : int  0 0 1 0 0 0 0 0 0 0 ...  
 $ S4      : int  0 0 0 1 0 0 0 0 0 0 ...  
 $ S5      : int  0 0 0 0 1 0 0 0 0 0 ...  
 $ S6      : int  0 0 0 0 0 1 0 0 0 0 ...  
 $ S7      : int  0 0 0 0 0 0 1 0 0 0 ...  
 $ S8      : int  0 0 0 0 0 0 0 1 0 0 ...  
 $ S9      : int  0 0 0 0 0 0 0 0 1 0 ...  
 $ S10     : int  0 0 0 0 0 0 0 0 0 1 ...  
 $ S11     : int  0 0 0 0 0 0 0 0 0 0 ...
```

In [81]:

```
# see some entries from begining ?  
head(data, 3)
```

t	Yt	Yt_trans	t.1	S1	S2	S3	S4	S5	S6	S7	S8	S9	S10	S11
1	501	4.731071	1	1	0	0	0	0	0	0	0	0	0	0
2	488	4.700077	2	0	1	0	0	0	0	0	0	0	0	0
3	504	4.738137	3	0	0	1	0	0	0	0	0	0	0	0

In [82]:

```
# model the data  
model = lm(Yt_trans ~ t.1+S1+S2+S3+S4+S5+S6+S7+S8+S9+S10+S11, data=data)
```

In [83]:

```
# summary of model?  
summary(model)
```

Call:

```
lm(formula = Yt_trans ~ t.1 + S1 + S2 + S3 + S4 + S5 + S6 + S7  
+  
    S8 + S9 + S10 + S11, data = data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.068082	-0.018755	0.001425	0.018808	0.079133

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	4.807e+00	8.463e-03	568.070	< 2e-16	***
t.1	3.515e-03	4.449e-05	79.009	< 2e-16	***
S1	-5.247e-02	1.055e-02	-4.971	1.75e-06	***
S2	-1.408e-01	1.055e-02	-13.342	< 2e-16	***
S3	-1.071e-01	1.055e-02	-10.151	< 2e-16	***
S4	4.988e-02	1.055e-02	4.728	5.05e-06	***
S5	2.542e-02	1.055e-02	2.410	0.0171	*
S6	1.902e-01	1.055e-02	18.031	< 2e-16	***
S7	3.825e-01	1.055e-02	36.266	< 2e-16	***
S8	4.134e-01	1.054e-02	39.201	< 2e-16	***
S9	7.142e-02	1.054e-02	6.773	2.47e-10	***
S10	5.064e-02	1.054e-02	4.803	3.66e-06	***
S11	-1.419e-01	1.054e-02	-13.463	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.0279 on 155 degrees of freedom

Multiple R-squared: 0.9884, Adjusted R-squared: 0.9875

F-statistic: 1098 on 12 and 155 DF, p-value: < 2.2e-16

In [84]:

```
# since y_trans is y^0.25  
data$predicted = (model$fitted.values)**4
```


In [85]:

```
head(data)
```

t	Yt	Yt_trans	t.1	S1	S2	S3	S4	S5	S6	S7	S8	S9	S10	S11	predic
1	501	4.731071	1	1	0	0	0	0	0	0	0	0	0	0	512.6
2	488	4.700077	2	0	1	0	0	0	0	0	0	0	0	0	477.0
3	504	4.738137	3	0	0	1	0	0	0	0	0	0	0	0	492.4
4	578	4.903227	4	0	0	0	1	0	0	0	0	0	0	0	563.0
5	545	4.831691	5	0	0	0	0	1	0	0	0	0	0	0	553.4
6	632	5.013942	6	0	0	0	0	0	1	0	0	0	0	0	634.3

In [86]:

```
ggplot(data, aes(t)) +  
  geom_line(aes(y = Yt, colour = "var0")) +  
  geom_line(aes(y = predicted, colour = "var1"))
```




In [87]:

```
# check the residuals (errors)
res = residuals(model)
```

In [88]:

```
plot(jitter(res)~jitter(Yt_trans), ylab="Residuals", xlab="Yt_trans", data=da,
abline(0,0))
```



Internal dummy

In [89]:

```
# let us use a file where R can create dummy internally
# file: t6-4 hotel_subset.csv

## Read file
data = read.csv('t6-4 hotel_subset.csv')

head(data, n=2)
#check the data types of all columns
sapply(data, class)
```

t	Month	Hotel.Occupancy..Y.	logY
1	Jan	501	2.700
2	Feb	488	2.688

```
t
'integer'
Month
'factor'
Hotel.Occupancy..Y.
'integer'
logY
'numeric'
```

In [90]:

```
data$newY = (data$Hotel.Occupancy..Y.)*0.25
model = lm(data$newY ~ t+Month, data=data)
summary(model)
```

Call:

```
lm(formula = data$newY ~ t + Month, data = data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.068082	-0.018755	0.001425	0.018808	0.079133

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	4.857e+00	8.300e-03	585.200	< 2e-16	***
t	3.515e-03	4.449e-05	79.009	< 2e-16	***
MonthAug	3.635e-01	1.054e-02	34.471	< 2e-16	***
MonthDec	-4.988e-02	1.055e-02	-4.728	5.05e-06	***
MonthFeb	-1.907e-01	1.054e-02	-18.084	< 2e-16	***
MonthJan	-1.023e-01	1.054e-02	-9.707	< 2e-16	***
MonthJul	3.326e-01	1.054e-02	31.541	< 2e-16	***
MonthJun	1.403e-01	1.054e-02	13.305	< 2e-16	***
MonthMar	-1.570e-01	1.054e-02	-14.889	< 2e-16	***
MonthMay	-2.446e-02	1.054e-02	-2.320	0.0216	*
MonthNov	-1.918e-01	1.055e-02	-18.186	< 2e-16	***
MonthOct	7.593e-04	1.055e-02	0.072	0.9427	
MonthSep	2.154e-02	1.055e-02	2.042	0.0428	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.0279 on 155 degrees of freedom

Multiple R-squared: 0.9884, Adjusted R-squared: 0.9875


F-statistic: 1098 on 12 and 155 DF, p-value: < 2.2e-16

In [91]:

```
# observe that the base level is not January but April as R chooses it intern
```



In [92]:

```
res <- residuals(model)
plot(jitter(res)~jitter(newY), ylab="Residuals", xlab="Y", data=data)
abline(0,0)
```



In [93]:

```
# since y_trans is y^0.25
data$predicted = (model$fitted.values)**4
ggplot(data, aes(t)) +
  geom_line(aes(y = Hotel.Occupancy..Y., colour = "var0")) +
  geom_line(aes(y = predicted, colour = "var1"))
```



B. Moving Averages

Let us use the data we saw during lecture for SP500 to perform a simple MA in R

Read in data

Data Import by R

R can import various data sources - txt, csv files and also has API to several data sources.

We illustrate with a csv data file.

The following reads in the data into a R variable 'dataframe' - named series.

In [94]:

```
series = read.csv('data_1995.csv')
```

In [95]:

```
# have a look at your data  
head(series, n =4)
```

Date	SP500	Dividend	Earnings	Consumer.Price.Index	Long.Interest.Rate
1/1/1995	465.25	13.18	31.25	150.3	7.78
1/2/1995	481.92	13.18	31.90	150.9	7.47
1/3/1995	493.15	13.17	32.55	151.4	7.20
1/4/1995	507.91	13.24	33.18	151.9	7.06

In [96]:

```
# extract the first 2 columns and discard the rest for now  
series = series[,1:2]
```

In [97]:

```
head(series)
```

Date	SP500
1/1/1995	465.25
1/2/1995	481.92
1/3/1995	493.15
1/4/1995	507.91
1/5/1995	523.81
1/6/1995	539.35

check if it is a valid time series object using function: `is.ts()`

- if not then we need to do some initial data processing

In [98]:

```
is.ts(series)
```

FALSE

Since it is not a valid TS object we need to change it to a TS object

In [99]:

```
sp500 = ts(series$SP500, frequency=12, start=c(1995, 1))
```

In [100]:

```
# how it looks now internally  
sp500
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	
1995	465.25	481.92	493.15	507.91	523.81	539.35	557.37	5
	59.11	578.77						
1996	614.42	649.54	647.07	647.17	661.23	668.50	644.07	6
	62.68	674.88						
1997	766.22	798.39	792.16	763.93	833.09	876.29	925.29	9
	27.24	937.02						
1998	963.36	1023.74	1076.83	1112.20	1108.42	1108.39	1156.58	10
	74.62	1020.64						
1999	1248.77	1246.58	1281.66	1334.76	1332.07	1322.55	1380.99	13
	27.49	1318.17						
2000	1425.59	1388.87	1442.21	1461.36	1418.48	1461.96	1473.00	14
	85.46	1468.05						
2001	1335.63	1305.75	1185.85	1189.84	1270.37	1238.71	1204.45	11
	78.50	1044.64						
2002	1140.21	1100.67	1153.79	1111.93	1079.25	1014.02	903.59	9
	12.55	867.81						
2003	895.84	837.03	846.63	890.03	935.96	988.00	992.54	9
	89.53	1019.44						
	Oct	Nov	Dec					
1995	582.92	595.53	614.57					
1996	701.46	735.67	743.25					
1997	951.16	938.92	962.37					
1998	1032.47	1144.43	1190.05					
1999	1300.01	1391.00	1428.68					
2000	1390.14	1378.04	1330.93					
2001	1076.59	1129.68	1144.93					
2002	854.63	909.93	899.18					
2003	1038.73	1049.90	1080.64					

In [101]:

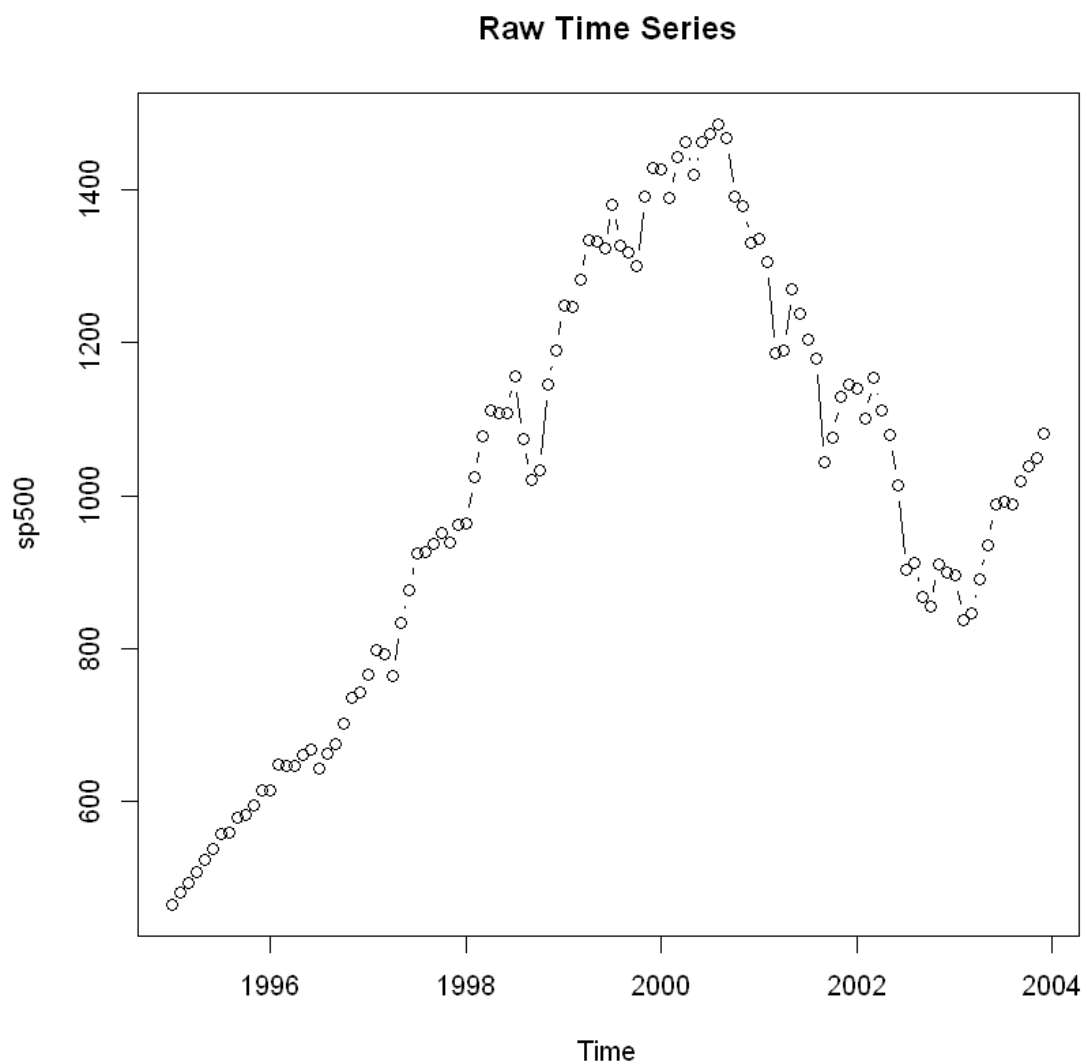
```
is.ts(sp500)
```

TRUE

plot your TS using ts.plot() command

In [102]:

```
ts.plot(sp500, main='Raw Time Series', type='b') # default type is 'l'
```



In [103]:

```
# Let us implement rolling mean
series2 = series %>%
  select(Date, SP500)%>%
  mutate(MA = rollmeanr(SP500, 2, fill = NA))
```

In [104]:

```
head(series2, n=4)
```

Date	SP500	MA
1/1/1995	465.25	NA
1/2/1995	481.92	473.585
1/3/1995	493.15	487.535
1/4/1995	507.91	500.530

In [105]:

```
is.ts(series2)
```

FALSE


In [106]:

```
#convert to a TS object  
sp500_ma = ts(series2$MA, frequency=12, start=c(1995, 2))
```

In [107]:

```
ts.plot(sp500,sp500_ma, lty=2:1, col=1:2, lwd=1:2, main="with moving average
```

with moving average superimposed



C. Forecating using AR and MA

Objectives of learning:

1. ARIMA (1,0,0) for AR
2. ARIMA (0,0,1) for MA
3. Sationarity test & ARIMA

In [108]:

```
## Read file "AmtrakBig.csv"  
series = read.csv('AmtrakBig.csv')  
head(series, n=4)
```

Month	Ridership	t	Season
Jan-91	1709	1	Jan
Feb-91	1621	2	Feb
Mar-91	1973	3	Mar
Apr-91	1812	4	Apr

In [110]:

```
#other useful commands
start(Rider)
end(Rider)

frequency(Rider)

#deltat(Rider)
cycle(Rider)

summary(Rider)
```

1991 1

2004 3


12

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1991	1	2	3	4	5	6	7	8	9	10	11	12
1992	1	2	3	4	5	6	7	8	9	10	11	12
1993	1	2	3	4	5	6	7	8	9	10	11	12
1994	1	2	3	4	5	6	7	8	9	10	11	12
1995	1	2	3	4	5	6	7	8	9	10	11	12
1996	1	2	3	4	5	6	7	8	9	10	11	12
1997	1	2	3	4	5	6	7	8	9	10	11	12
1998	1	2	3	4	5	6	7	8	9	10	11	12
1999	1	2	3	4	5	6	7	8	9	10	11	12
2000	1	2	3	4	5	6	7	8	9	10	11	12
2001	1	2	3	4	5	6	7	8	9	10	11	12
2002	1	2	3	4	5	6	7	8	9	10	11	12
2003	1	2	3	4	5	6	7	8	9	10	11	12
2004	1	2	3									

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1361	1698	1831	1822	1967	2223

In [111]:

```
ts.plot(Rider)
```



In []:

AR(1)

In [112]:

```
# AR using ARIMA

arima_100 = arima(Rider, order=c(1,0,0))
arima_100
```

Call:

```
arima(x = Rider, order = c(1, 0, 0))
```

Coefficients:

	ar1	intercept
	0.5680	1823.8221
s.e.	0.0656	27.1868

sigma^2 estimated as 22285: log likelihood = -1021.73, aic = 2049.47

In [113]:


```
# forecast the values
forecasted_values = forecast(arima_100, 12)
forecasted_values
accuracy(forecasted_values)
plot(forecasted_values, lwd=2)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2004	1998.880	1807.567	2190.193	1706.292	2291.468
May 2004	1923.263	1703.238	2143.287	1586.765	2259.760
Jun 2004	1880.308	1651.788	2108.828	1530.817	2229.800
Jul 2004	1855.909	1624.714	2087.104	1502.327	2209.491
Aug 2004	1842.049	1609.997	2074.100	1487.157	2196.940
Sep 2004	1834.176	1601.849	2066.503	1478.862	2189.489
Oct 2004	1829.703	1597.287	2062.119	1474.254	2185.153
Nov 2004	1827.163	1594.718	2059.607	1471.670	2182.656
Dec 2004	1825.720	1593.266	2058.174	1470.212	2181.227
Jan 2005	1824.900	1592.443	2057.357	1469.388	2180.412
Feb 2005	1824.434	1591.977	2056.892	1468.921	2179.948
Mar 2005	1824.170	1591.712	2056.628	1468.656	2179.684

	ME	RMSE	MAE	MPE	MAPE	MASE	AC
Training set	0.5395504	149.2823	120.06	-0.6812322	6.784501	1.472945	-0.02920



Forecasts from ARIMA(1,0,0) with non-zero mean



In []:

MA(1)

In [114]:

```
# AR using ARIMA

arima_001 = arima(Rider, order=c(0,0,1))
arima_001
```

Call:

```
arima(x = Rider, order = c(0, 0, 1))
```

Coefficients:

	ma1	intercept
	0.5067	1823.2098
s.e.	0.0686	18.5777

sigma^2 estimated as 24273: log likelihood = -1028.48, aic = 2062.96

In [115]:

```
# forecast the values
forecasted_values = forecast(arima_001, 50)
forecasted_values
accuracy(forecasted_values)
plot(forecasted_values, lwd=2)
```

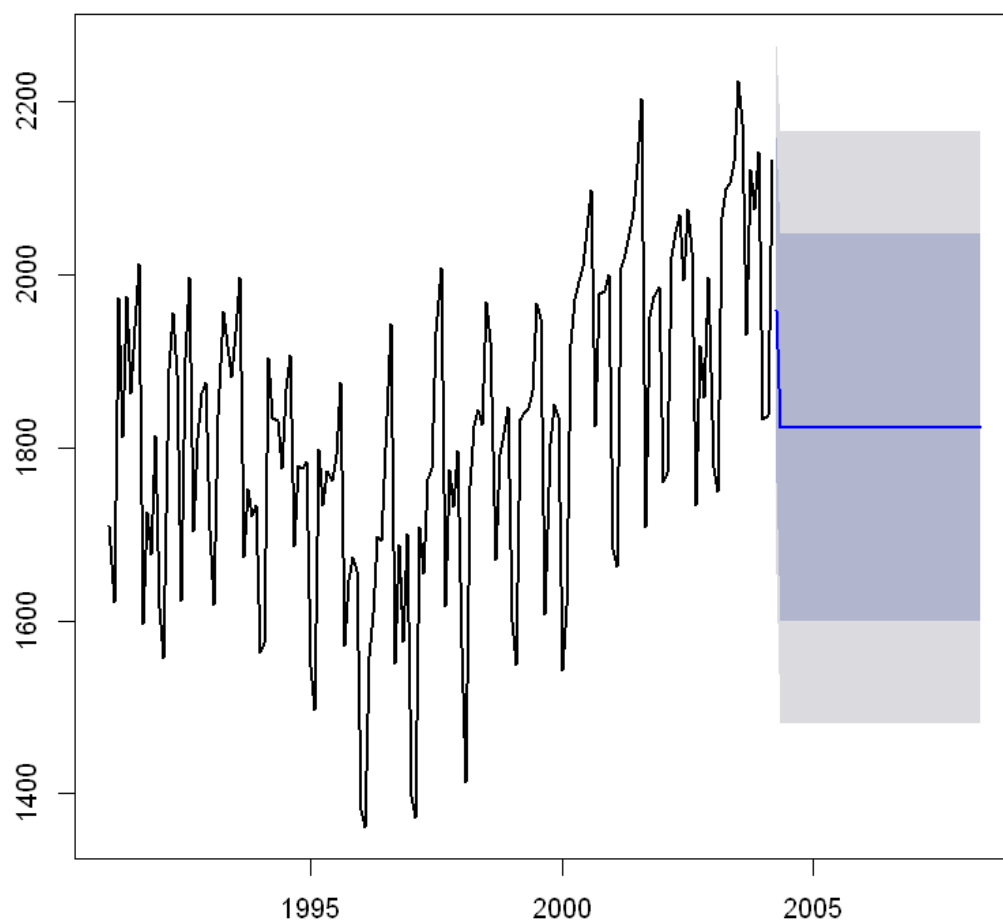
	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2004	1959.529	1759.866	2159.192	1654.171	2264.887
May 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Jun 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Jul 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Aug 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Sep 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Oct 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Nov 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Dec 2004	1823.210	1599.379	2047.041	1480.889	2165.530
Jan 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Feb 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Mar 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Apr 2005	1823.210	1599.379	2047.041	1480.889	2165.530
May 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Jun 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Jul 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Aug 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Sep 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Oct 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Nov 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Dec 2005	1823.210	1599.379	2047.041	1480.889	2165.530
Jan 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Feb 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Mar 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Apr 2006	1823.210	1599.379	2047.041	1480.889	2165.530
May 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Jun 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Jul 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Aug 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Sep 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Oct 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Nov 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Dec 2006	1823.210	1599.379	2047.041	1480.889	2165.530
Jan 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Feb 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Mar 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Apr 2007	1823.210	1599.379	2047.041	1480.889	2165.530
May 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Jun 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Jul 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Aug 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Sep 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Oct 2007	1823.210	1599.379	2047.041	1480.889	2165.530

Nov 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Dec 2007	1823.210	1599.379	2047.041	1480.889	2165.530
Jan 2008	1823.210	1599.379	2047.041	1480.889	2165.530
Feb 2008	1823.210	1599.379	2047.041	1480.889	2165.530
Mar 2008	1823.210	1599.379	2047.041	1480.889	2165.530
Apr 2008	1823.210	1599.379	2047.041	1480.889	2165.530
May 2008	1823.210	1599.379	2047.041	1480.889	2165.530

	ME	RMSE	MAE	MPE	MAPE	MASE	
Training set	-0.04981039	155.7977	128.715	-0.8087708	7.256321	1.579127	0.096



Forecasts from ARIMA(0,0,1) with non-zero mean



In []:

Self Study: Additional Section

When AR or MA does not give desirable results

Try more sophisticated technique : ARIMA(p, d, q)

- ARIMA: autoregressive integrated moving average
- Key point: we need a stationary time series
- Differencing: method to remove trend. first order if we take difference from the previous value.
- Apply ARIMA on a stationarised time series
- Use ACF and PACF curves to find p and q
- p : for AR (PACF chart)
- d : for differencing order
 - subtract the trend from the series using differencing with the previous values
- q : for MA (ACF chart)

Is it a stationary TS?

- Let us test using Augmented Dickey-Fuller Test
what is my null hypothesis?
- it is a non-stionary time series

In [116]:

```
adfTest(Rider)
```

Title:

Augmented Dickey-Fuller Test

Test Results:

PARAMETER:

Lag Order: 1

STATISTIC:

Dickey-Fuller: -0.1713

P VALUE:

0.5598

Description:

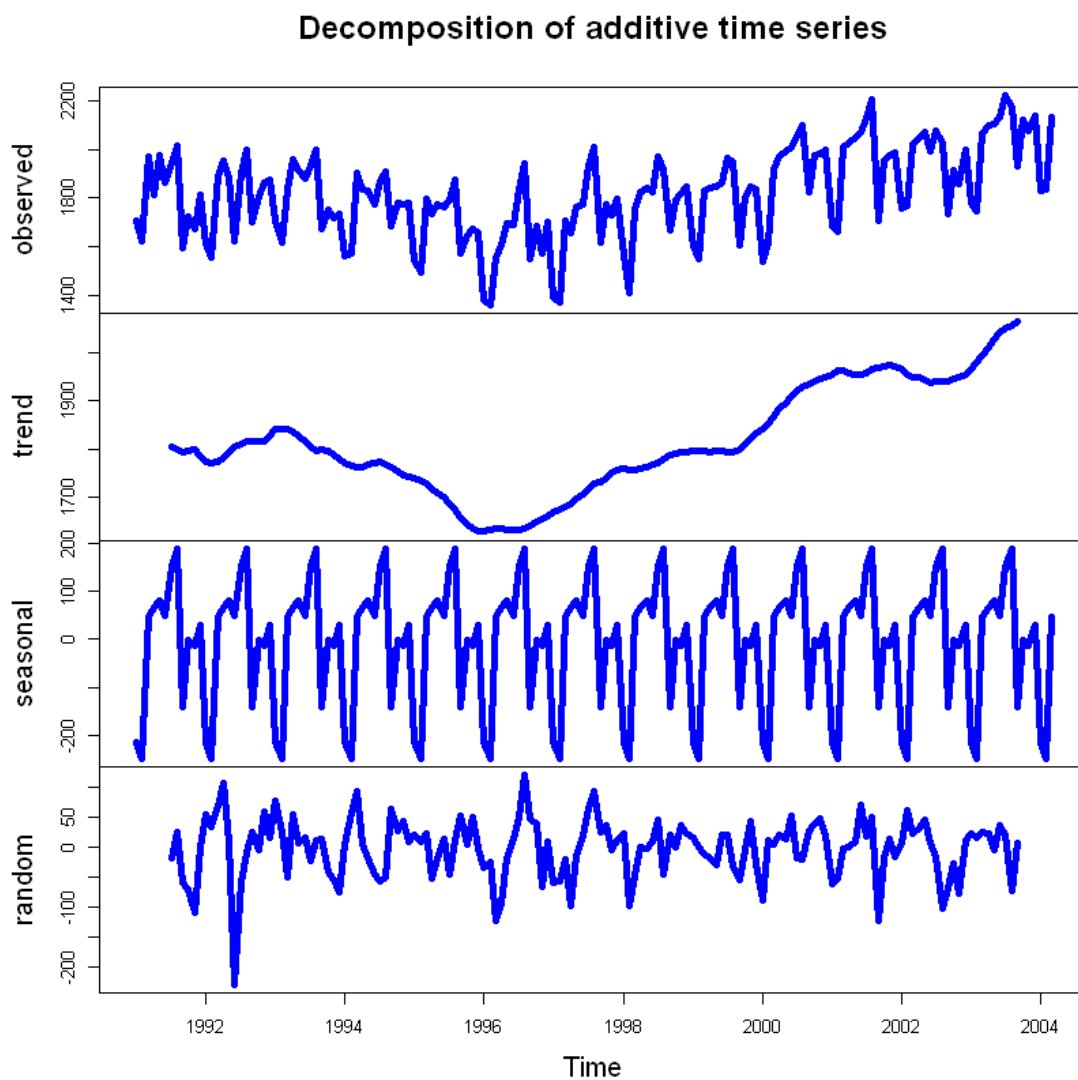
Thu May 10 10:18:00 2018 by user: isspcs

Observations ?

- looking at the p-value we cannot reject the null hypothesis

In [117]:

```
# visualize the various components of TS
AmCom = decompose(Rider)
plot(AmCom, lwd=4, col="blue")
```




In [118]:

```
# as the TS appears non-stationary
# perform differencing of order 1
# to achieve stationarity
# using the "diff()" function in R
```

In [119]:

```
Rider_diff = diff(Rider, differences = 1)
plot.ts(Rider_diff,col="blue")
```



In [120]:

```
# test again the differenced TS  
adfTest(Rider_diff)
```

Warning message in adfTest(Rider_diff):
"p-value smaller than printed p-value"

Title:

Augmented Dickey-Fuller Test

Test Results:

PARAMETER:

Lag Order: 1

STATISTIC:

Dickey-Fuller: -13.8351

P VALUE:

0.01

Description:


Thu May 10 10:18:10 2018 by user: isspcs

View auto correlation of time series data


ACF and PACF to discern p, q at d=1

In [121]:

```
Acf(Rider_diff, 104) # for p  
Pacf(Rider_diff, 104)# for q
```



Series Rider_diff



In [122]:

```
arima_111 = arima(Rider, order=c(1,1,1))
arima_111

# forecast the values
forecasted_values<-forecast(arima_111, 12)
forecasted_values
accuracy(forecasted_values)
plot(forecasted_values, lwd=2)
```

Call:

```
arima(x = Rider, order = c(1, 1, 1))
```

Coefficients:

```
          ar1      ma1
      0.3488 -0.9163
s.e.  0.0832  0.0294
```

sigma^2 estimated as 20714: log likelihood = -1009.93, aic = 2025.87

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2004	2054.032	1869.586	2238.478	1771.946	2336.118
May 2004	2026.836	1825.878	2227.795	1719.497	2334.176
Jun 2004	2017.350	1811.787	2222.914	1702.969	2331.732
Jul 2004	2014.042	1806.224	2221.860	1696.211	2331.872
Aug 2004	2012.887	1803.438	2222.336	1692.563	2333.212
Sep 2004	2012.485	1801.603	2223.366	1689.969	2335.000
Oct 2004	2012.344	1800.102	2224.587	1687.747	2336.942
Nov 2004	2012.295	1798.721	2225.869	1685.662	2338.929
Dec 2004	2012.278	1797.389	2227.168	1683.633	2340.924
Jan 2005	2012.272	1796.077	2228.468	1681.631	2342.914
Feb 2005	2012.270	1794.779	2229.762	1679.645	2344.895
Mar 2005	2012.270	1793.489	2231.050	1677.674	2346.865


	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	11.64368	143.4707	120.6574	0.003759092	6.787888	1.480274

◀

▶

▲

Forecasts from ARIMA(1,1,1)



Improve Further using AutoArima

- R has an improved autoarima method

In [123]:

```
fitAutoArima = auto.arima(Rider)
fitAutoArima

#acf(residuals(fitAutoArima), 24, lwd=2)
Box.test(residuals(fitAutoArima), lag=12, type="Ljung")

f3<-forecast(fitAutoArima, 12)
f3
accuracy(f3)
plot(f3, lwd=2)
```

Series: Rider
ARIMA(2,1,1)(2,1,2)[12]

Coefficients:

	ar1	ar2	ma1	sar1	sar2	sma1	sma2
	0.4372	0.1093	-0.8426	0.7065	-0.1721	-1.4182	0.5751
s.e.	0.1463	0.1014	0.1006	0.3406	0.1708	0.3417	0.2574

sigma^2 estimated as 3710: log likelihood=-808.88
AIC=1633.76 AICc=1634.81 BIC=1657.63


Box-Ljung test

data: residuals(fitAutoArima)
X-squared = 4.8051, df = 12, p-value = 0.9642

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2004	2163.965	2085.894	2242.036	2044.566	2283.364
May 2004	2186.394	2095.567	2277.222	2047.485	2325.304
Jun 2004	2201.974	2102.276	2301.672	2049.499	2354.449
Jul 2004	2288.869	2183.096	2394.642	2127.103	2450.635
Aug 2004	2286.234	2175.662	2396.807	2117.128	2455.340
Sep 2004	1969.172	1854.540	2083.803	1793.858	2144.486
Oct 2004	2154.715	2036.465	2272.964	1973.868	2335.562
Nov 2004	2142.919	2021.334	2264.503	1956.972	2328.866
Dec 2004	2187.248	2062.520	2311.976	1996.493	2378.002
Jan 2005	1906.287	1778.554	2034.020	1710.937	2101.637
Feb 2005	1900.663	1770.033	2031.294	1700.882	2100.445
Mar 2005	2201.317	2067.874	2334.760	1997.233	2405.401

	ME	RMSE	MAE	MPE	MAPE	MASE	
Training set	6.020336	56.95005	43.41794	0.2691171	2.425085	0.5326688	-0.0150

Forecasts from ARIMA(2,1,1)(2,1,2)[12]



Type *Markdown* and LaTeX: α^2

Some more additional implementations

Holt Winter's Additive model

In [130]:

```
fitHW <- HoltWinters(Rider, seasonal="additive")  
  
## Check the model summary and performance  
  
fitHW
```

Holt-Winters exponential smoothing with trend and additive seasonal component.

Call:

```
HoltWinters(x = Rider, seasonal = "additive")
```

Smoothing parameters:

```
alpha: 0.3291363  
beta : 0.005983517  
gamma: 0.7483261
```

Coefficients:

```
      [,1]  
a 2024.1666702  
b   0.7482538  
s1 122.4331329  
s2 123.2120521  
s3 113.9806811  
s4 188.9536220  
s5 157.6268160
```

In [131]:

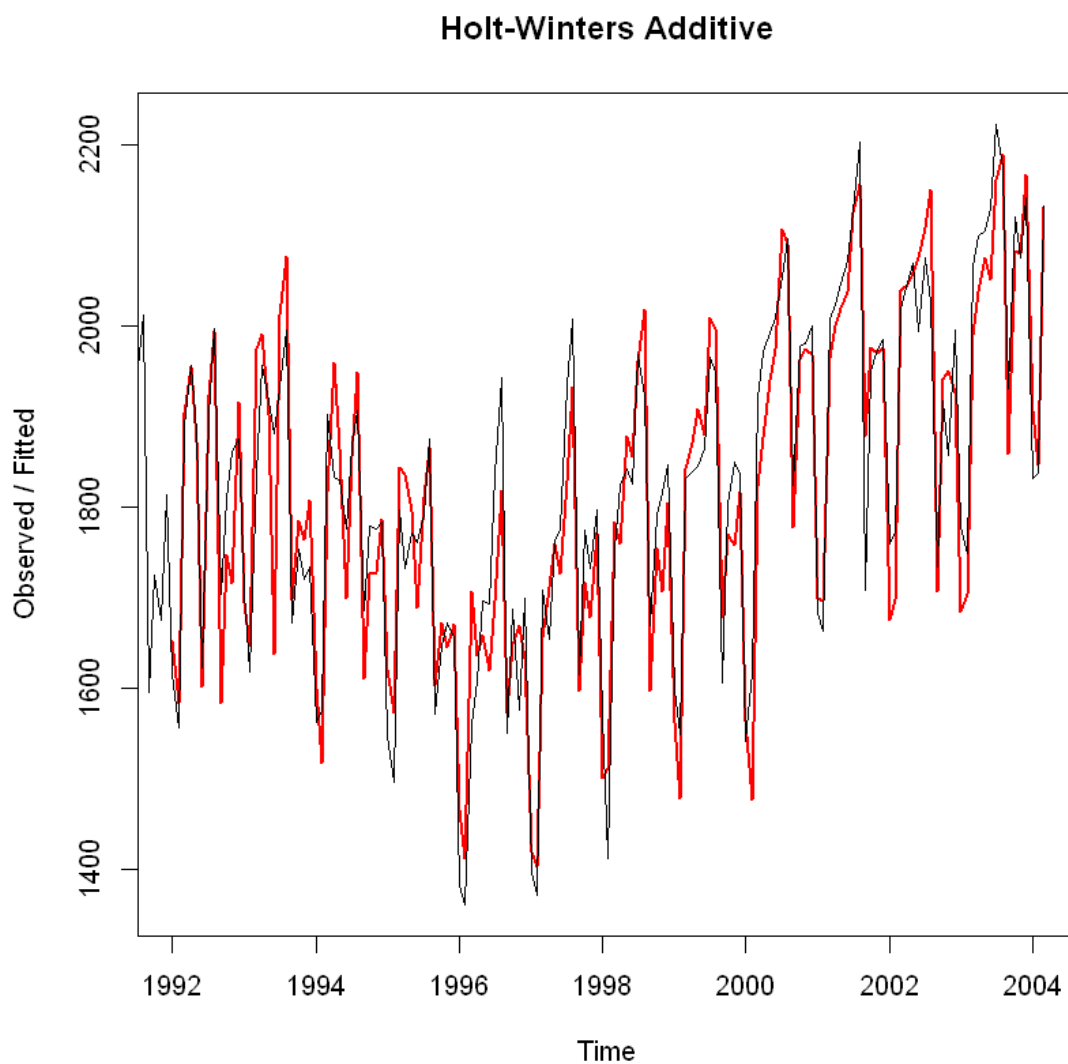
```
Box.test(residuals(fitHW), lag=12, type="Ljung")

## >>> Plot the fitted Time Series

plot(fitHW, main="Holt-Winters Additive",lwd=2)
```

Box-Ljung test

data: residuals(fitHW)
X-squared = 24.08, df = 12, p-value = 0.01984



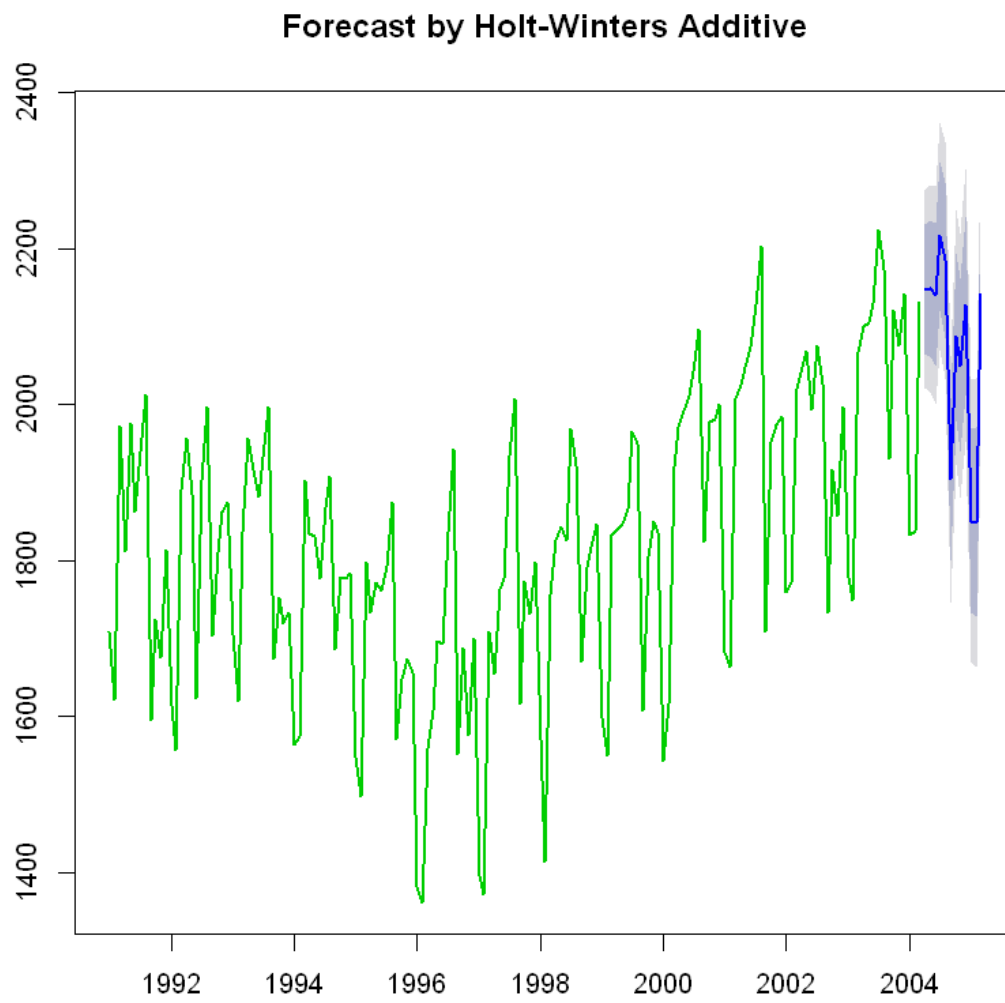
In [125]:

```
## Make forecast & check performance of model
f <-forecast(fitHW,12)
f
accuracy(f)
plot(f, lwd=2, col="green3", main="Forecast by Holt-Winters Additive")
```

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2004		2147.348	2064.724	2229.972	2020.985	2273.711
May 2004		2148.875	2061.840	2235.911	2015.766	2281.984
Jun 2004		2140.392	2049.109	2231.675	2000.787	2279.997
Jul 2004		2216.113	2120.725	2311.501	2070.230	2361.997
Aug 2004		2185.535	2086.166	2284.904	2033.563	2337.507
Sep 2004		1903.559	1800.318	2006.800	1745.666	2061.452
Oct 2004		2086.845	1979.829	2193.860	1923.179	2250.511
Nov 2004		2049.811	1939.108	2160.513	1880.506	2219.116
Dec 2004		2127.278	2012.967	2241.590	1952.454	2302.103
Jan 2005		1850.867	1733.018	1968.717	1670.632	2031.102
Feb 2005		1848.752	1727.430	1970.075	1663.205	2034.299
Mar 2005		2140.857	2016.121	2265.594	1950.089	2331.626

	ME	RMSE	MAE	MPE	MAPE	MASE	AC
Training set	6.746409	64.60549	51.18466	0.2987725	2.865657	0.6279541	0.23280





Fit a Holt Winter's Multiplicative model

In [126]:

```
fitHW1 <- HoltWinters(Rider, seasonal="multiplicative")
fitHW1

Box.test(residuals(fitHW1), lag=12, type="Ljung")
plot(fitHW1, main="Holt-Winters Multiplicative", lwd=2)
```

Holt-Winters exponential smoothing with trend and multiplicative seasonal component.

Call:

```
HoltWinters(x = Rider, seasonal = "multiplicative")
```

Smoothing parameters:

```
alpha: 0.3089791
beta : 0.002171449
gamma: 0.7283249
```

Coefficients:


```
      [,1]
a 2023.8929691
b  -0.3953454
s1  1.0618703
s2  1.0629445
s3  1.0585834
s4  1.0978202
s5  1.0833999
s6  0.9379491
s7  1.0305579
s8  1.0121733
s9  1.0517197
s10 0.9109065
s11 0.9071362
s12 1.0535303
```

Box-Ljung test

data: residuals(fitHW1)

X-squared = 29.849, df = 12, p-value = 0.002942

Holt-Winters Multiplicative



In [127]:

```
f1 <-forecast(fitHW1, 12)
f1
accuracy(f1)
plot(f1, lwd=2, col="green3", main="Forecast by Holt-Winters Multiplicative")
```

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95	
Apr 2004		2148.692	2079.909	2217.474	2043.498	2253.886	
May 2004		2150.445	2076.825	2224.066	2037.853	2263.038	
Jun 2004		2141.204	2063.107	2219.301	2021.764	2260.644	
Jul 2004		2220.135	2136.778	2303.491	2092.651	2347.618	
Aug 2004		2190.544	2103.541	2277.546	2057.485	2323.603	
Sep 2004		1896.084	1810.159	1982.008	1764.674	2027.493	
Oct 2004		2082.887	1989.205	2176.569	1939.612	2226.161	
Nov 2004		2045.329	1948.823	2141.835	1897.736	2192.922	
Dec 2004		2124.826	2022.639	2227.013	1968.545	2281.107	
Jan 2005		1839.976	1742.255	1937.697	1690.525	1989.427	
Feb 2005		1832.002	1730.947	1933.056	1677.452	1986.551	
Mar 2005		2127.234	1922.421	2332.048	1813.999	2440.470	
	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	7.096806	66.23209	52.73617	0.3245168	2.950346	0.6469885	0.2610309

In []:

Fit an Exponential Smoothing Model

In [128]:

```
fitExp<-ets(Rider)#holt's
fitExp

Box.test(residuals(fitExp), lag=12, type="Ljung")

f2<-forecast(fitExp, 12)
f2
accuracy(f2)
plot(f2, lwd=2, col="green3")
```

ETS(A,N,A)

Call:

```
ets(y = Rider)
```

Smoothing parameters:

alpha = 0.5625

gamma = 1e-04

Initial states:

l = 1813.1983

s=27.3613 -13.4603 1.85 -139.0846 189.3363 152.2884

47.3404 85.0458 58.763 47.9488 -249.1373 -208.2517

sigma: 58.216

	AIC	AICc	BIC
	2113.703	2117.060	2159.737

Box-Ljung test

data: residuals(fitExp)


X-squared = 10.65, df = 12, p-value = 0.5591

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2004	2140.417	2065.811	2215.024	2026.316	2254.518
May 2004	2166.705	2081.104	2252.306	2035.790	2297.620
Jun 2004	2129.000	2033.664	2224.335	1983.197	2274.802
Jul 2004	2233.944	2129.780	2338.108	2074.639	2393.249
Aug 2004	2270.993	2158.692	2383.294	2099.244	2442.742
Sep 2004	1942.575	1822.688	2062.461	1759.224	2125.925
Oct 2004	2083.511	1956.491	2210.531	1889.250	2277.771
Nov 2004	2068.201	1934.427	2201.975	1863.612	2272.790
Dec 2004	2109.026	1968.824	2249.228	1894.605	2323.447
Jan 2005	1873.405	1727.057	2019.754	1649.584	2097.227
Feb 2005	1832.528	1680.281	1984.776	1599.686	2065.371
Mar 2005	2129.613	1971.685	2287.541	1888.083	2371.143

ME	RMSE	MAE	MPE	MAPE	MASE
----	------	-----	-----	------	------

	ME	RMSE	MAE	MPE	MAPE	MASE	
Training set	3.001498	55.59394	43.62935	0.08490419	2.436975	0.5352624	0.0430

Forecasts from ETS(A,N,A)



More Illustrations

Decomposition

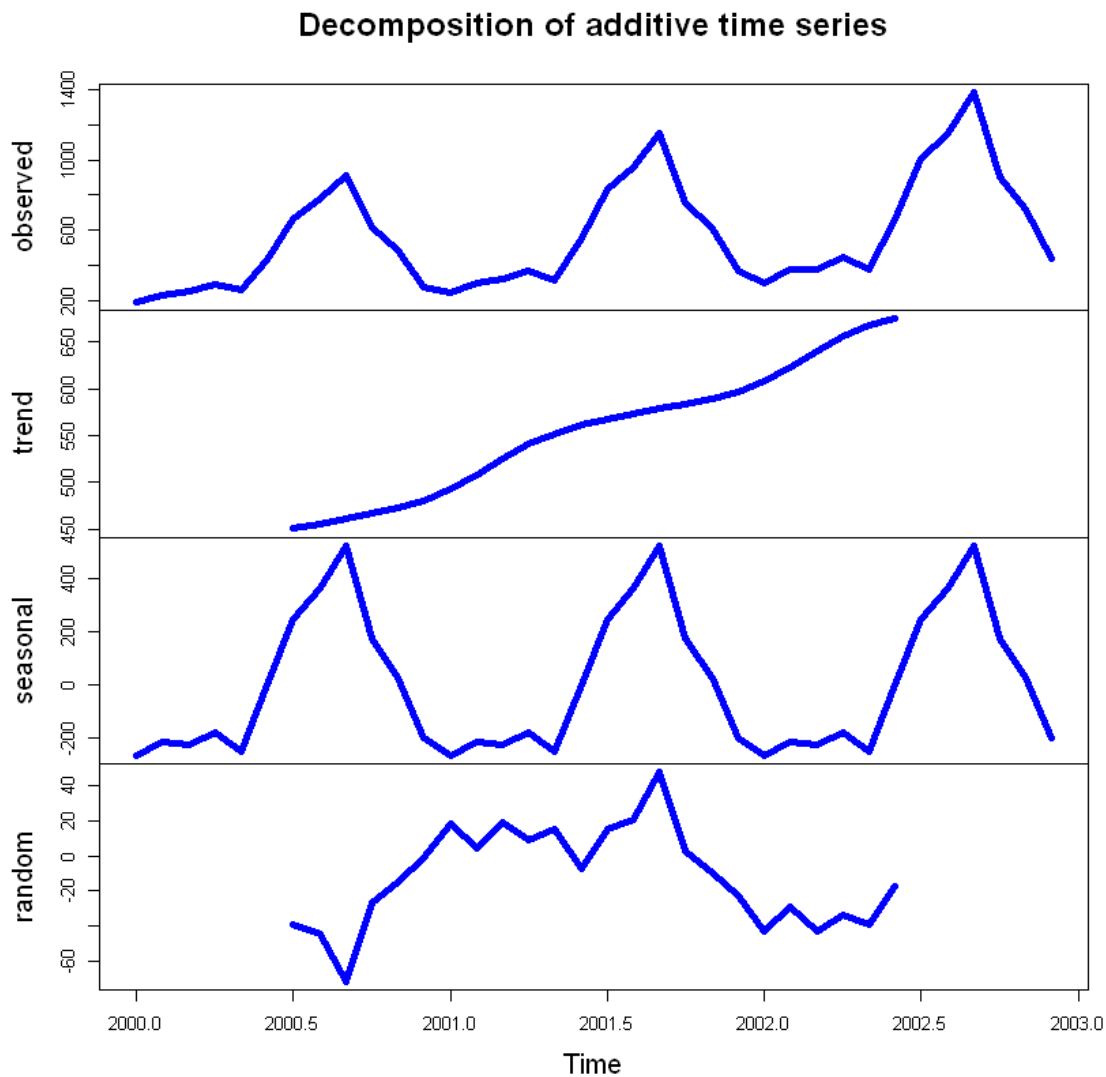
```
## Read file "cola-data.csv"
cola_series = read.csv('Cola_data.csv')
head(cola_series, n=4)
```

```
cola_ts = ts(cola_series$Sales..Y., frequency=12, start=c(2000, 1))
cola_ts
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2000	189	229	249	289	260	431	660	777	915	613	485	277
2001	244	296	319	370	313	556	831	960	1152	759	607	371
2002	298	378	373	443	374	660	1004	1153	1388	904	715	441

In [8]:

```
# visualize the various components of TS  
cola_com = decompose(cola_ts)  
plot(cola_com, lwd=4, col="blue")
```



Arima

In [9]:

```
## Read file "Electricity raw data.csv"
elec_series = read.csv('Electricity raw data.csv')
head(elec_series, n=4)
```

Year	Month	Electricity
1	2005-01	448.1
1	2005-02	437.4
1	2005-03	480.0
1	2005-04	533.9

In [10]:


```
elec_ts = ts(elec_series$Electricity, frequency=12, start=c(2005, 1))
elec_ts
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2005	448.1	437.4	480.0	533.9	535.3	560.6	537.7	529.1	516.7	520.4	518.9	478.3
2006	465.9	452.7	470.4	511.8	522.4	533.8	513.8	544.9	542.0	526.4	535.8	495.6
2007	474.2	445.0	452.4	511.7	543.5	556.8	563.9	555.8	514.4	515.5	548.7	490.7
2008	467.9	473.6	450.2	492.3	544.8	573.5	526.7	510.9	520.0	523.0	536.4	484.9
2009	477.2	461.3	462.0	511.4	557.5	610.5	593.6	573.1	555.8	558.3	563.9	516.9
2010	491.5	504.8	548.5	586.0	574.5	607.0	576.2	540.6	540.1	558.7	582.4	537.1
2011	500.2	462.6	466.6	525.3	566.0	603.8	587.0	591.6	583.6	558.7	542.0	506.8
2012	504.2	505.9	509.5	541.8	558.1	608.9	595.8	586.5	574.2	561.1	569.1	525.8
2013	514.0	498.9	499.1	569.2	602.7	630.6	612.9	618.6	561.6	555.8	571.7	531.0

In [12]:

```
# visualize the various components of TS
elec_com = decompose(elec_ts)
plot(elec_com, lwd=4, col="blue")
```

Decomposition of additive time series



In [13]:

```
arima_111 = arima(elec_ts, order=c(1,1,1))
arima_111

# forecast the values
forecasted_values<-forecast(arima_111, 12)
forecasted_values
accuracy(forecasted_values)
```

Call:

```
arima(x = elec_ts, order = c(1, 1, 1))
```

Coefficients:

```
          ar1      ma1
      0.1040  0.1880
s.e.  0.2943  0.2874
```

```
sigma^2 estimated as 778.4:  log likelihood = -508.03,  aic = 1
022.06
```


	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 2014		518.2974	482.5428	554.0521	463.6154	572.9794
Feb 2014		516.9761	458.5608	575.3914	427.6376	606.3147
Mar 2014		516.8387	441.6867	591.9907	401.9036	631.7737
Apr 2014		516.8244	427.9762	605.6725	380.9429	652.7058
May 2014		516.8229	416.1190	617.5267	362.8096	670.8362
Jun 2014		516.8227	405.5183	628.1271	346.5973	687.0481
Jul 2014		516.8227	395.8430	637.8024	331.8003	701.8451
Aug 2014		516.8227	386.8862	646.7592	318.1020	715.5434
Sep 2014		516.8227	378.5082	655.1372	305.2889	728.3565
Oct 2014		516.8227	370.6095	663.0359	293.2089	740.4365
Nov 2014		516.8227	363.1161	670.5293	281.7487	751.8967
Dec 2014		516.8227	355.9714	677.6740	270.8219	762.8235

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	0.4864186	27.77006	22.0688	0.008373021	4.151499	1.024718



In [14]:

```
plot(forecasted_values, lwd=2)
```



Auto-Arima

In [11]:

```
fitAutoArima = auto.arima(elec_ts)
fitAutoArima

#acf(residuals(fitAutoArima), 24, lwd=2)
Box.test(residuals(fitAutoArima), lag=12, type="Ljung")

f3<-forecast(fitAutoArima, 12)
f3
accuracy(f3)
plot(f3, lwd=2)
```

Series: elec_ts
ARIMA(3,0,2)(2,1,0)[12] with drift

Coefficients:

	ar1	ar2	ar3	ma1	ma2	sar1	sar
2 drift							
	-0.3843	-0.2009	0.3264	1.4106	0.8780	-0.7045	-0.210
8 0.5925							
s.e.	0.1205	0.1120	0.1321	0.0765	0.0733	0.1153	0.112
5 0.1737							

sigma^2 estimated as 213.8: log likelihood=-393.38
AIC=804.77 AICc=806.86 BIC=827.85


Box-Ljung test

data: residuals(fitAutoArima)
X-squared = 8.121, df = 12, p-value = 0.7756

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 2014	515.8311	497.0920	534.5702	487.1722	544.4900
Feb 2014	516.7737	489.9229	543.6246	475.7089	557.8386
Mar 2014	508.0988	480.7302	535.4675	466.2422	549.9555
Apr 2014	558.1334	530.7639	585.5028	516.2754	599.9914
May 2014	590.6335	562.7875	618.4794	548.0468	633.2202
Jun 2014	625.7235	597.8761	653.5709	583.1346	668.3124
Jul 2014	611.9949	584.1346	639.8553	569.3862	654.6037
Aug 2014	612.6727	584.7364	640.6089	569.9478	655.3975
Sep 2014	584.7389	556.7935	612.6842	542.0001	627.4776
Oct 2014	572.5568	544.6083	600.5053	529.8133	615.3003
Nov 2014	578.7264	550.7610	606.6919	535.9569	621.4959
Dec 2014	536.1626	508.1923	564.1330	493.3856	578.9397

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.09804067	13.19901	9.106145	-0.08277868	1.719975	0.4228246

Forecasts from ARIMA(3,0,2)(2,1,0)[12] with drift



In []: